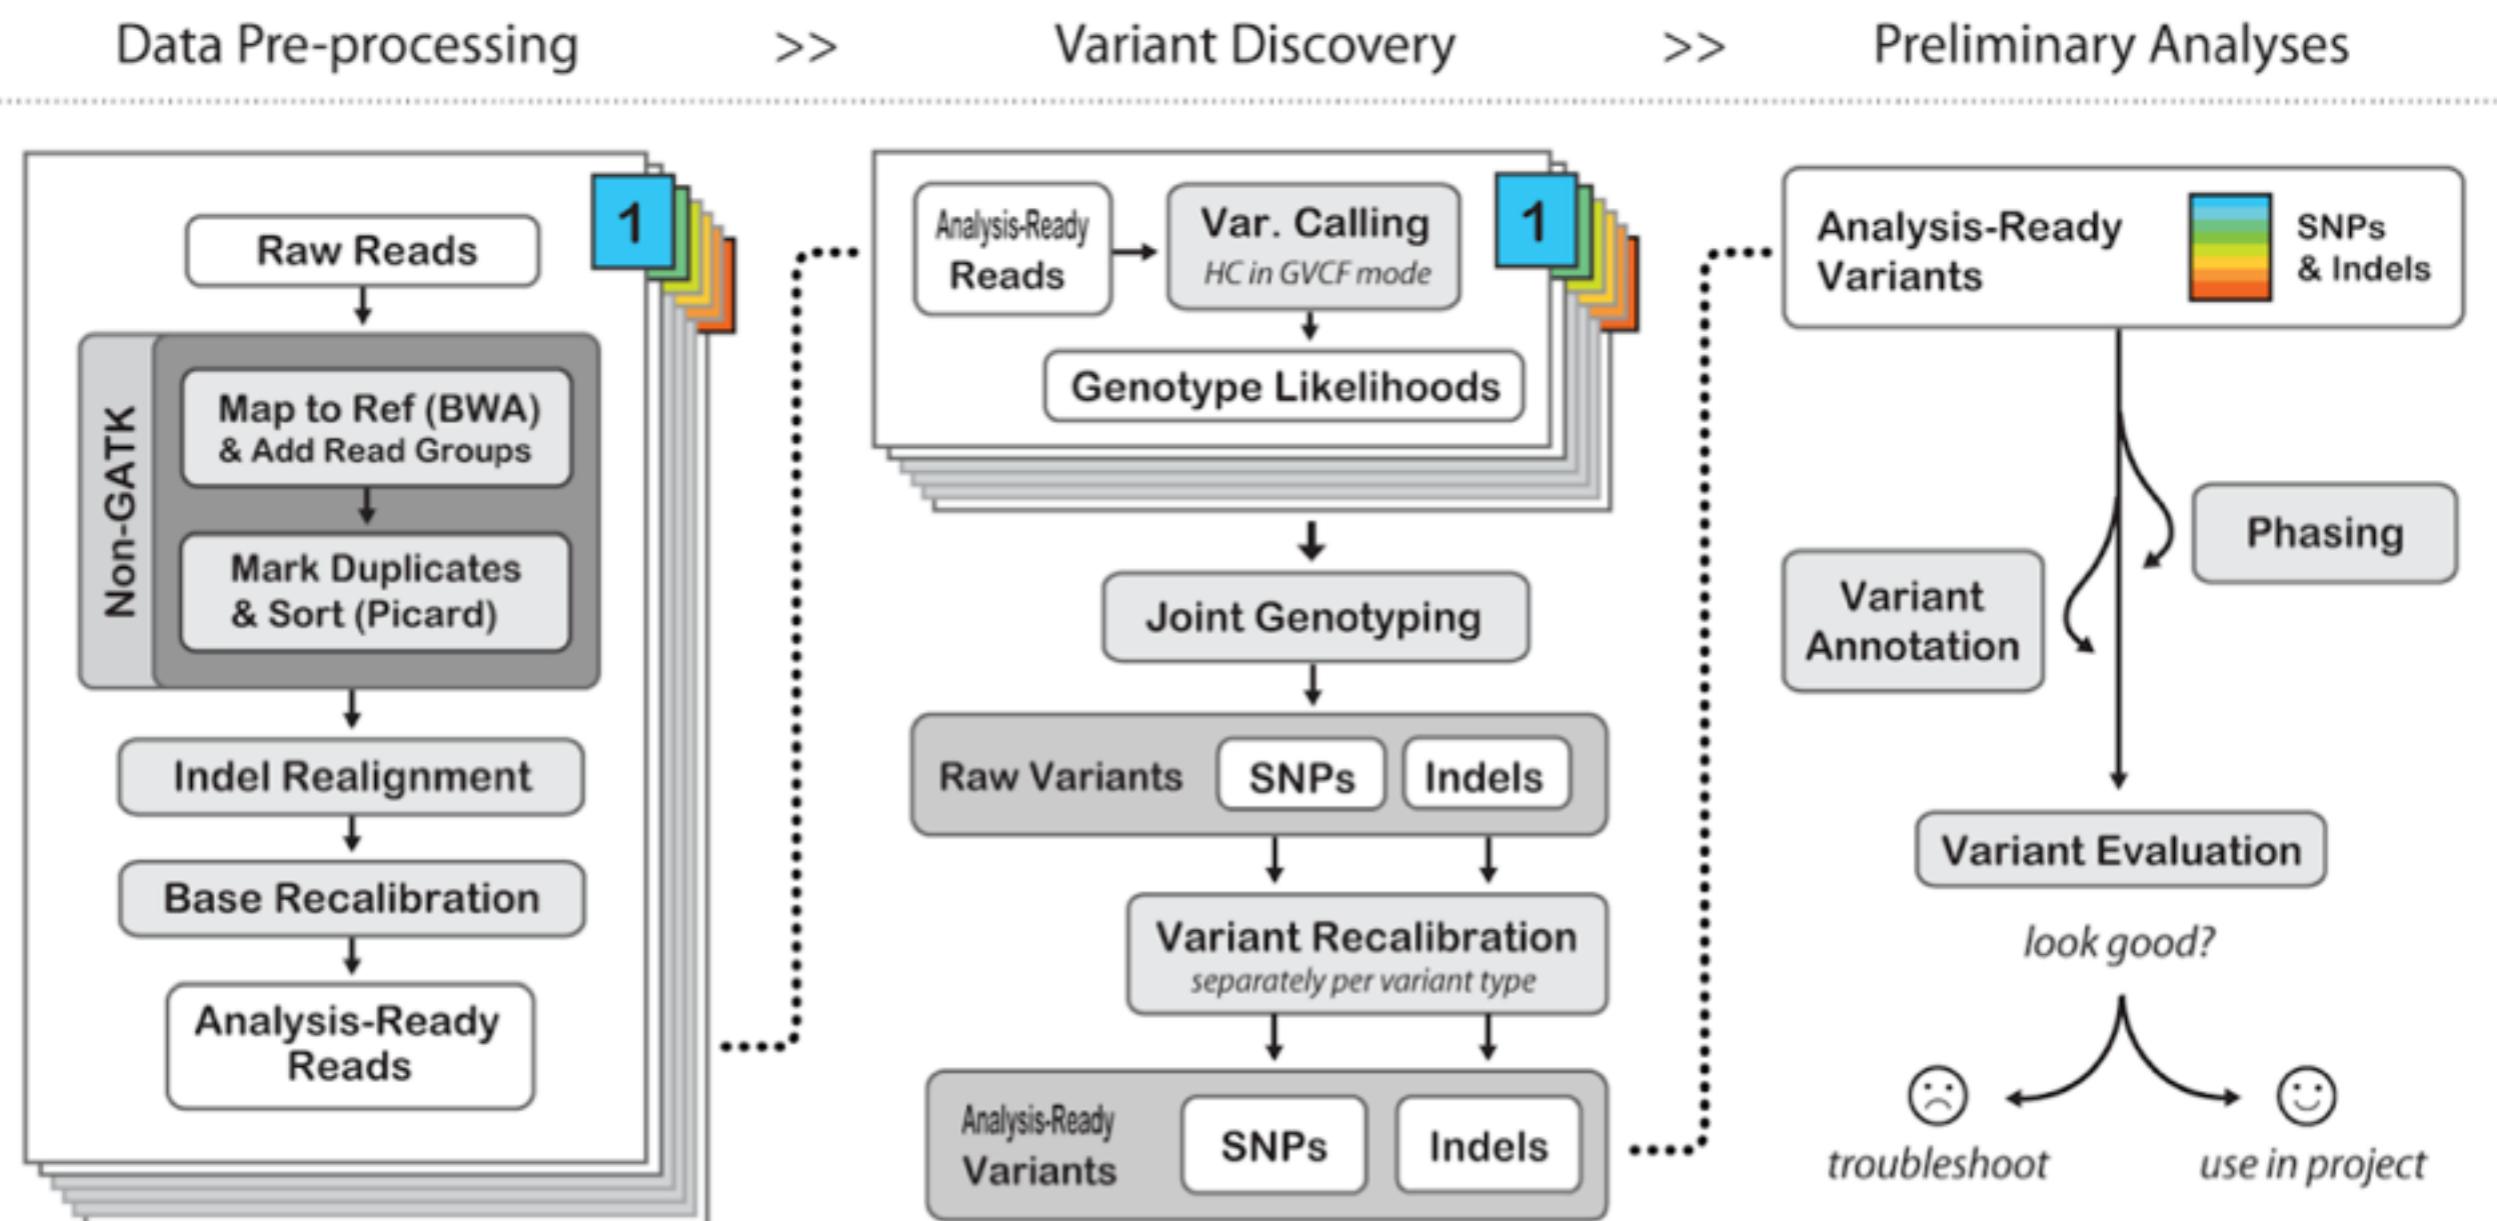


Topic 6: Variant calling using GATK

Learning Goals

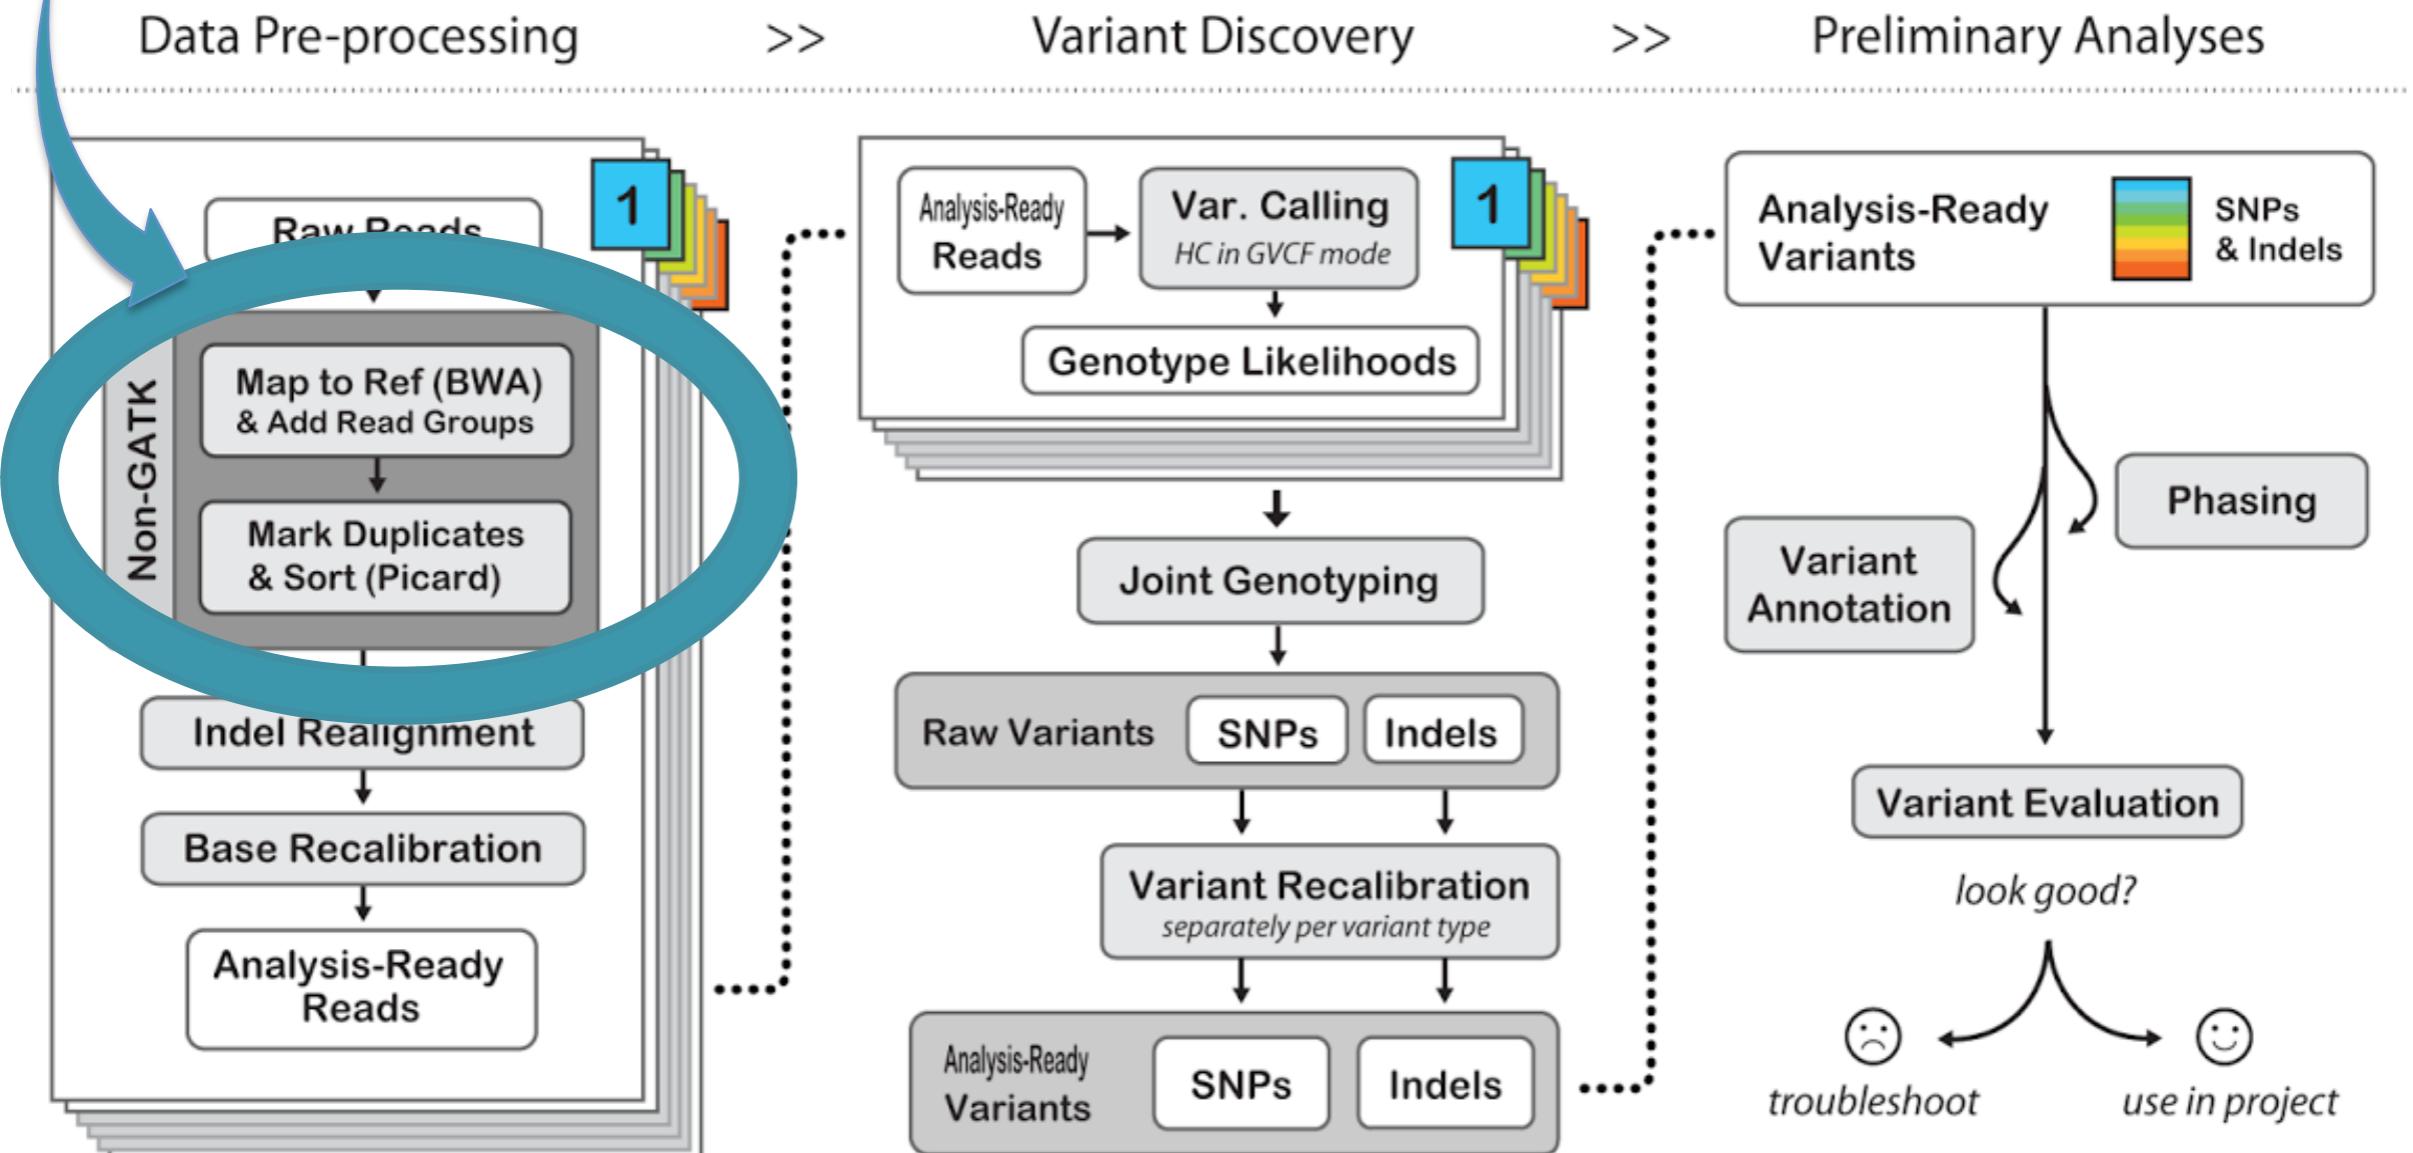
- Define the steps involved in SNP calling and what they are doing.
- Understand the reason for haplotype based SNP calling.
- Define the N+1 problem in genotyping.

GATK Best Practises for Variant Discovery in DNAsed



We are here in the Best Practices workflow

Mapping and Marking Duplicates



BAM headers: an essential part of a BAM file

```
@HD VN:1.0 GO:none SO:coordinate  
@SQ SN:chrM LN:16571  
@SQ SN:chr1 LN:247249719  
@SQ SN:chr2 LN:242951149  
[cut for clarity]  
@SQ SN:chr9 LN:140273252  
@SQ SN:chr10 LN:135374737  
@SQ SN:chr11 LN:134452384  
[cut for clarity]  
@SQ SN:chr22 LN:49691432  
@SQ SN:chrX LN:154913754  
@SQ SN:chrY LN:57772954  
  
@RG ID:20FUK.1 PL:illumina PU:20FUKAAXX100202.1 LB:Solexa-18483 SM:NA12878 CN:BI  
@RG ID:20FUK.2 PL:illumina PU:20FUKAAXX100202.2 LB:Solexa-18484 SM:NA12878 CN:BI  
@RG ID:20FUK.3 PL:illumina PU:20FUKAAXX100202.3 LB:Solexa-18483 SM:NA12878 CN:BI  
@RG ID:20FUK.4 PL:illumina PU:20FUKAAXX100202.4 LB:Solexa-18484 SM:NA12878 CN:BI  
@RG ID:20FUK.5 PL:illumina PU:20FUKAAXX100202.5 LB:Solexa-18483 SM:NA12878 CN:BI  
@RG ID:20FUK.6 PL:illumina PU:20FUKAAXX100202.6 LB:Solexa-18484 SM:NA12878 CN:BI  
@RG ID:20FUK.7 PL:illumina PU:20FUKAAXX100202.7 LB:Solexa-18483 SM:NA12878 CN:BI  
@RG ID:20FUK.8 PL:illumina PU:20FUKAAXX100202.8 LB:Solexa-18484 SM:NA12878 CN:BI  
  
@PG ID:BWA VN:0.5.7 CL:tk  
@PG ID:GATK TableRecalibration VN:1.0.2864
```

Required: Standard header

Essential: contigs of aligned reference sequence. Should be in karyotypic order.

Essential: read groups. Carries platform (PL), library (LB), and sample (SM) information. Each read is associated with a read group

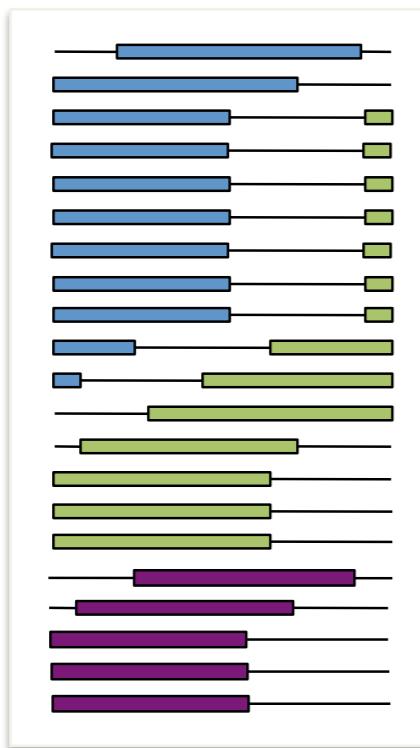
20FUKAAXX100202:1:1:12730:189900 163 chrM 1 60 101M = 282 381
GATCACAGGTCTATCACCCTATTAAACCACTCACGGGAGCTCCATGCATTGGTA...[more bases]
?BA@A>BBBBACBBAC@BBCBBCBC@BC@CAC@:BBCBBCACAACBABCBCAB...[more quals]
RG:Z:20FUK.1 NM:i:1 SM:i:37 AM:i:37 MD:Z:72G28 MQ:i:60 PG:Z:BWA UQ:i:33

Useful: Data processing tools applied to the reads

Official specification in <http://samtools.sourceforge.net/SAM1.pdf>

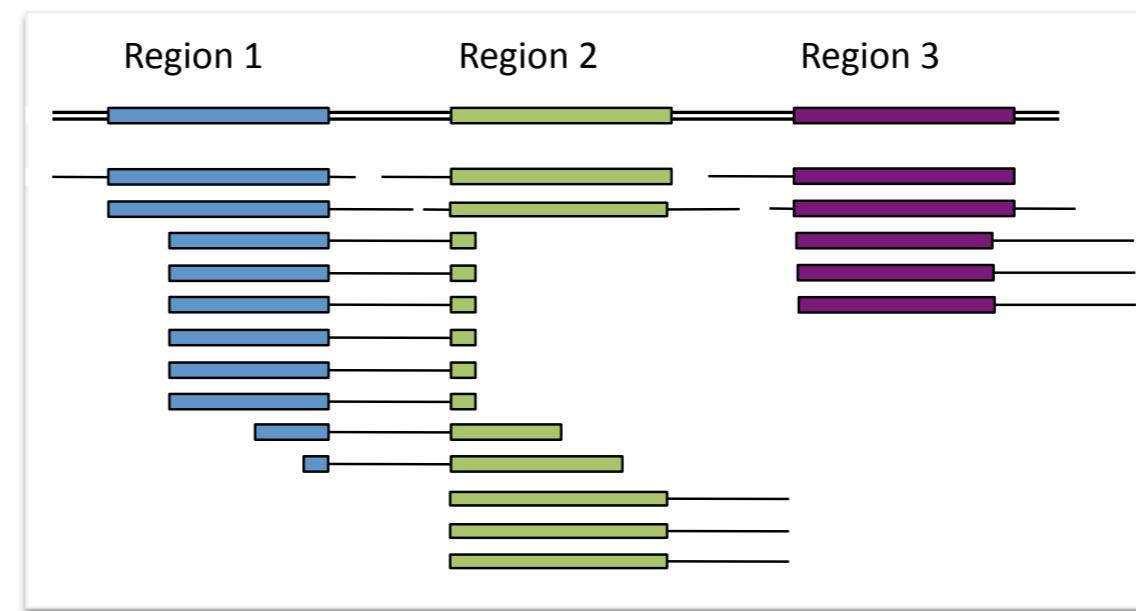
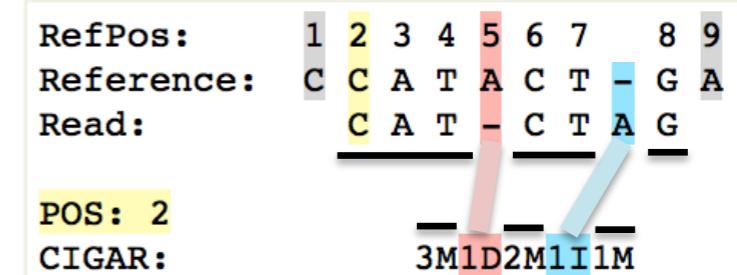
Mapping short reads to a reference is simple in principle

Enormous pile of short reads from NGS



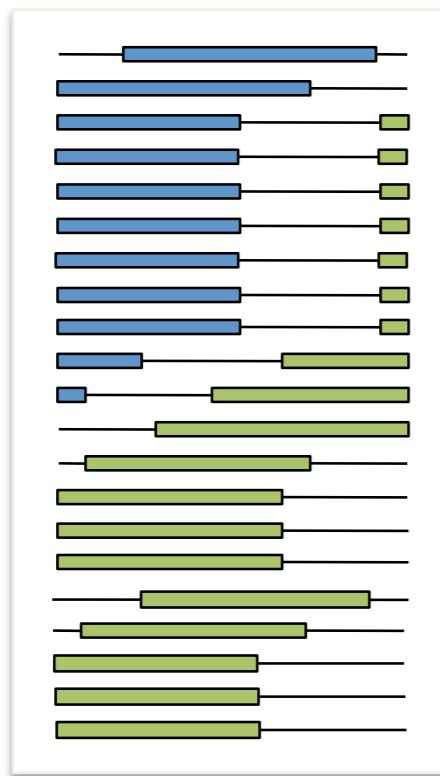
Mapping and alignment algorithms

Identify where the read matches the reference sequence and record match details as CIGAR string



But mapping is actually very hard because of mismatches (true mutations or sequencing errors), duplicated regions etc.

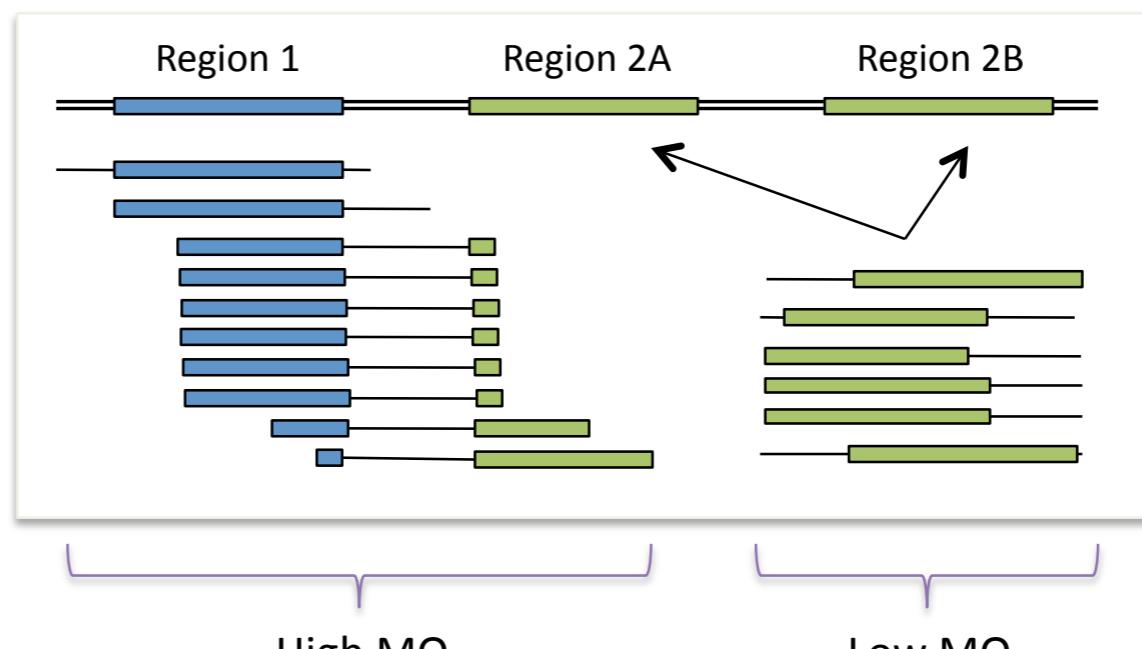
Enormous pile of short reads from NGS



Mapping and alignment algorithms

Mapping algorithms account for this by choosing the most likely placement

→ mapping quality (MQ)

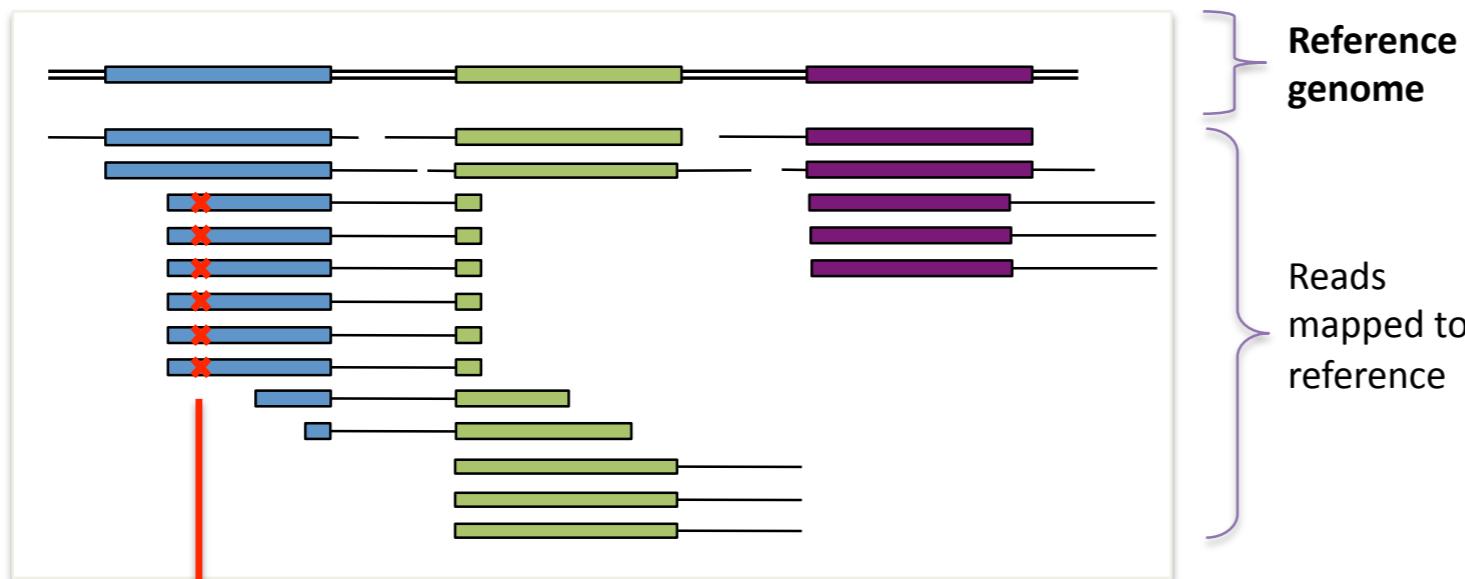


For more information see:

Li and Homer (2010). A survey of sequence alignment algorithms for next-generation sequencing.
Briefings in Bioinformatics.

The reason why duplicates are bad

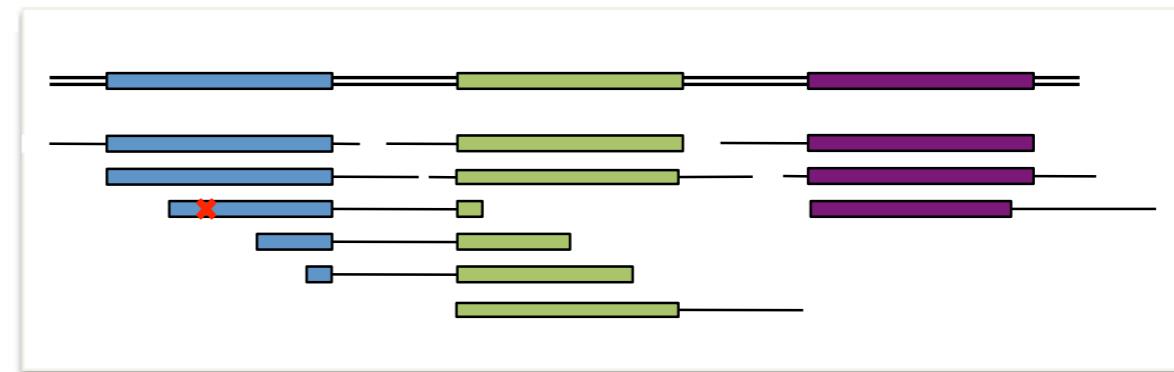
✖ = sequencing error propagated in duplicates



FP variant call
(bad)



After marking duplicates, the GATK will only see :



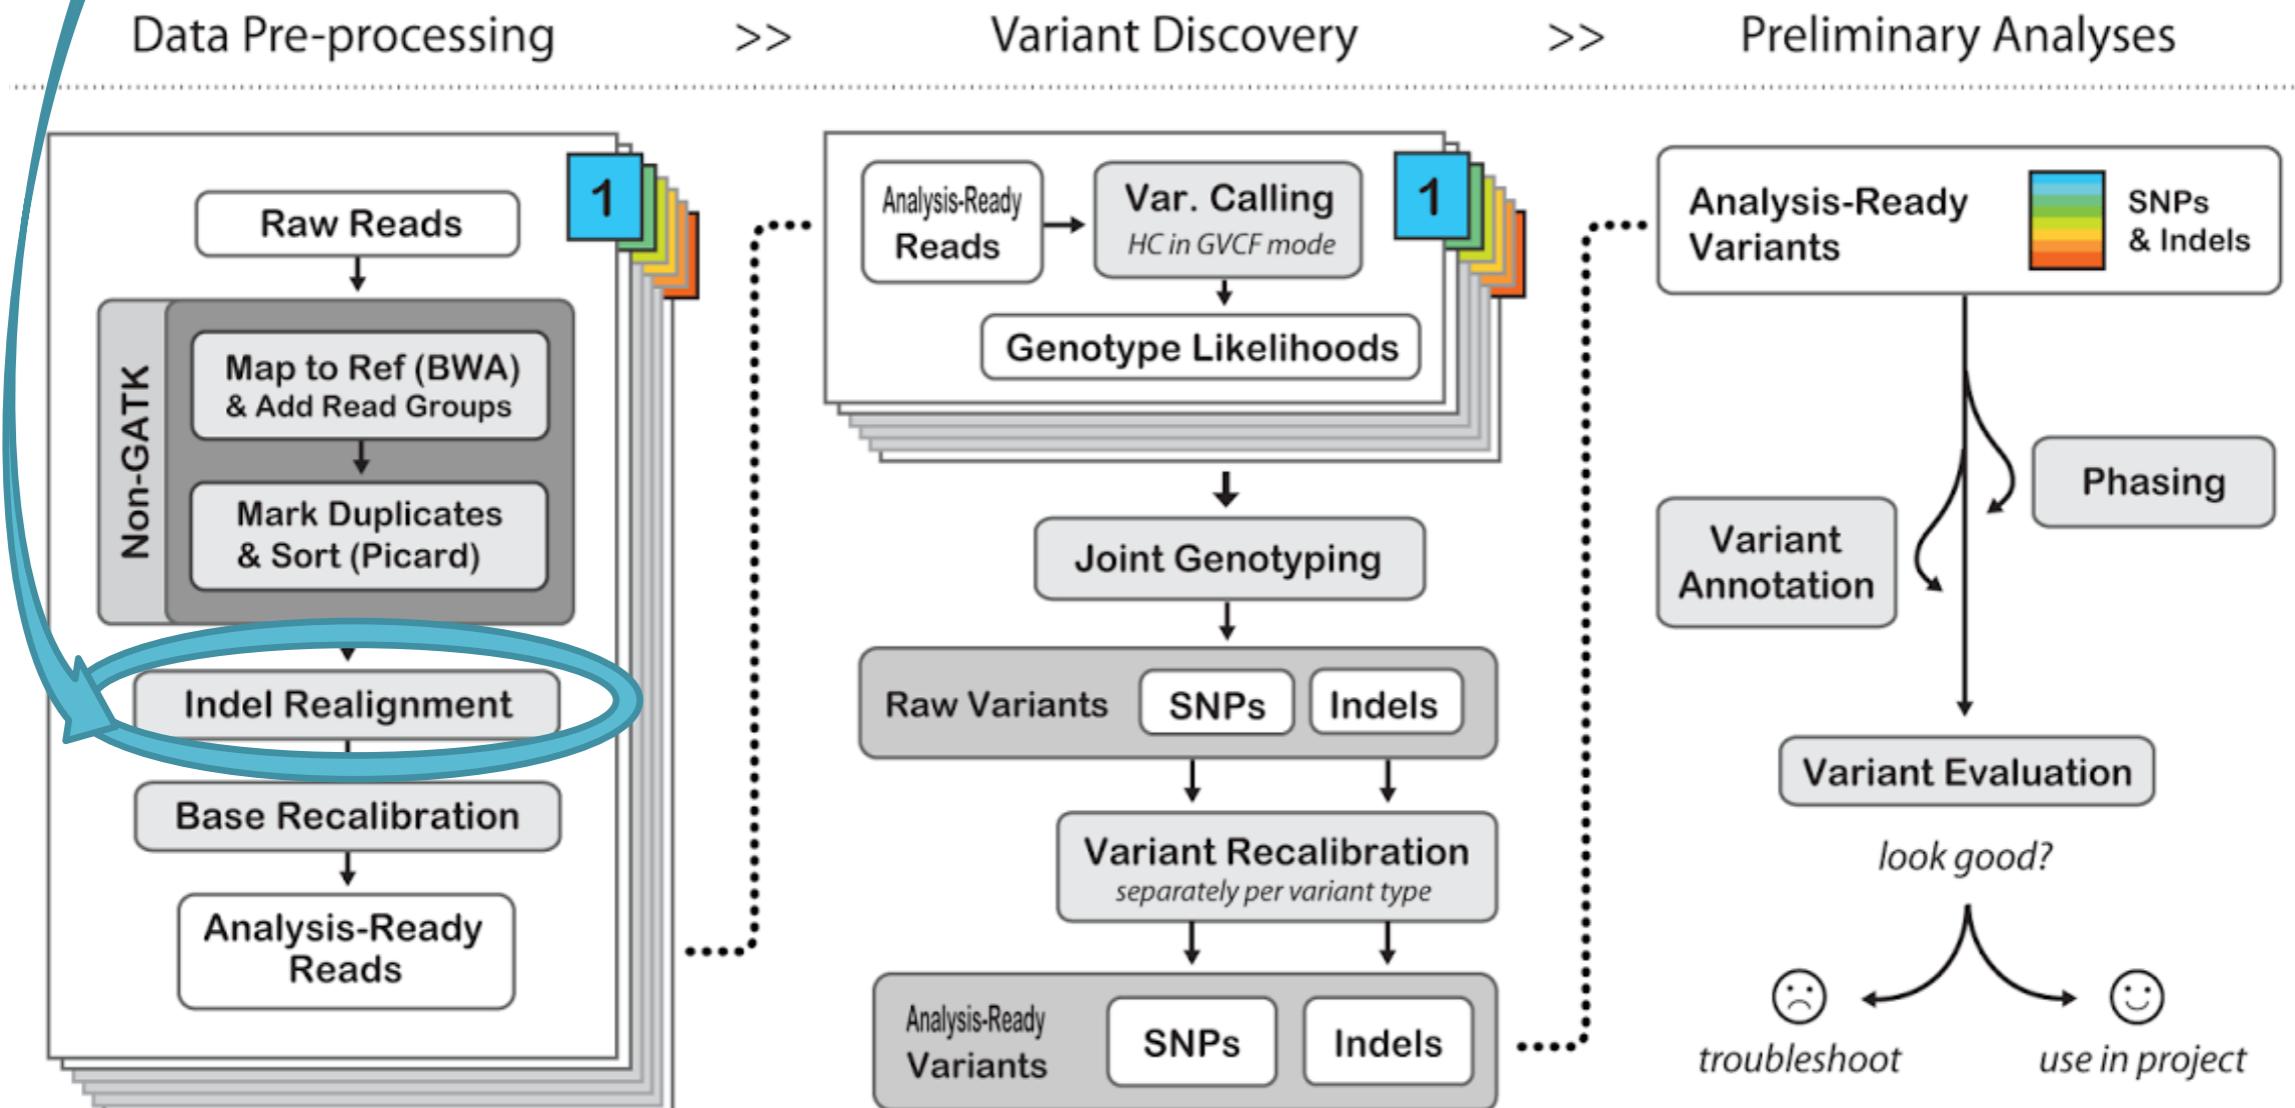
... and thus be more likely to make the right call

Easy to identify: duplicate reads have the same starting position and same CIGAR string



We are here in the Best Practices workflow

Indel Realignment

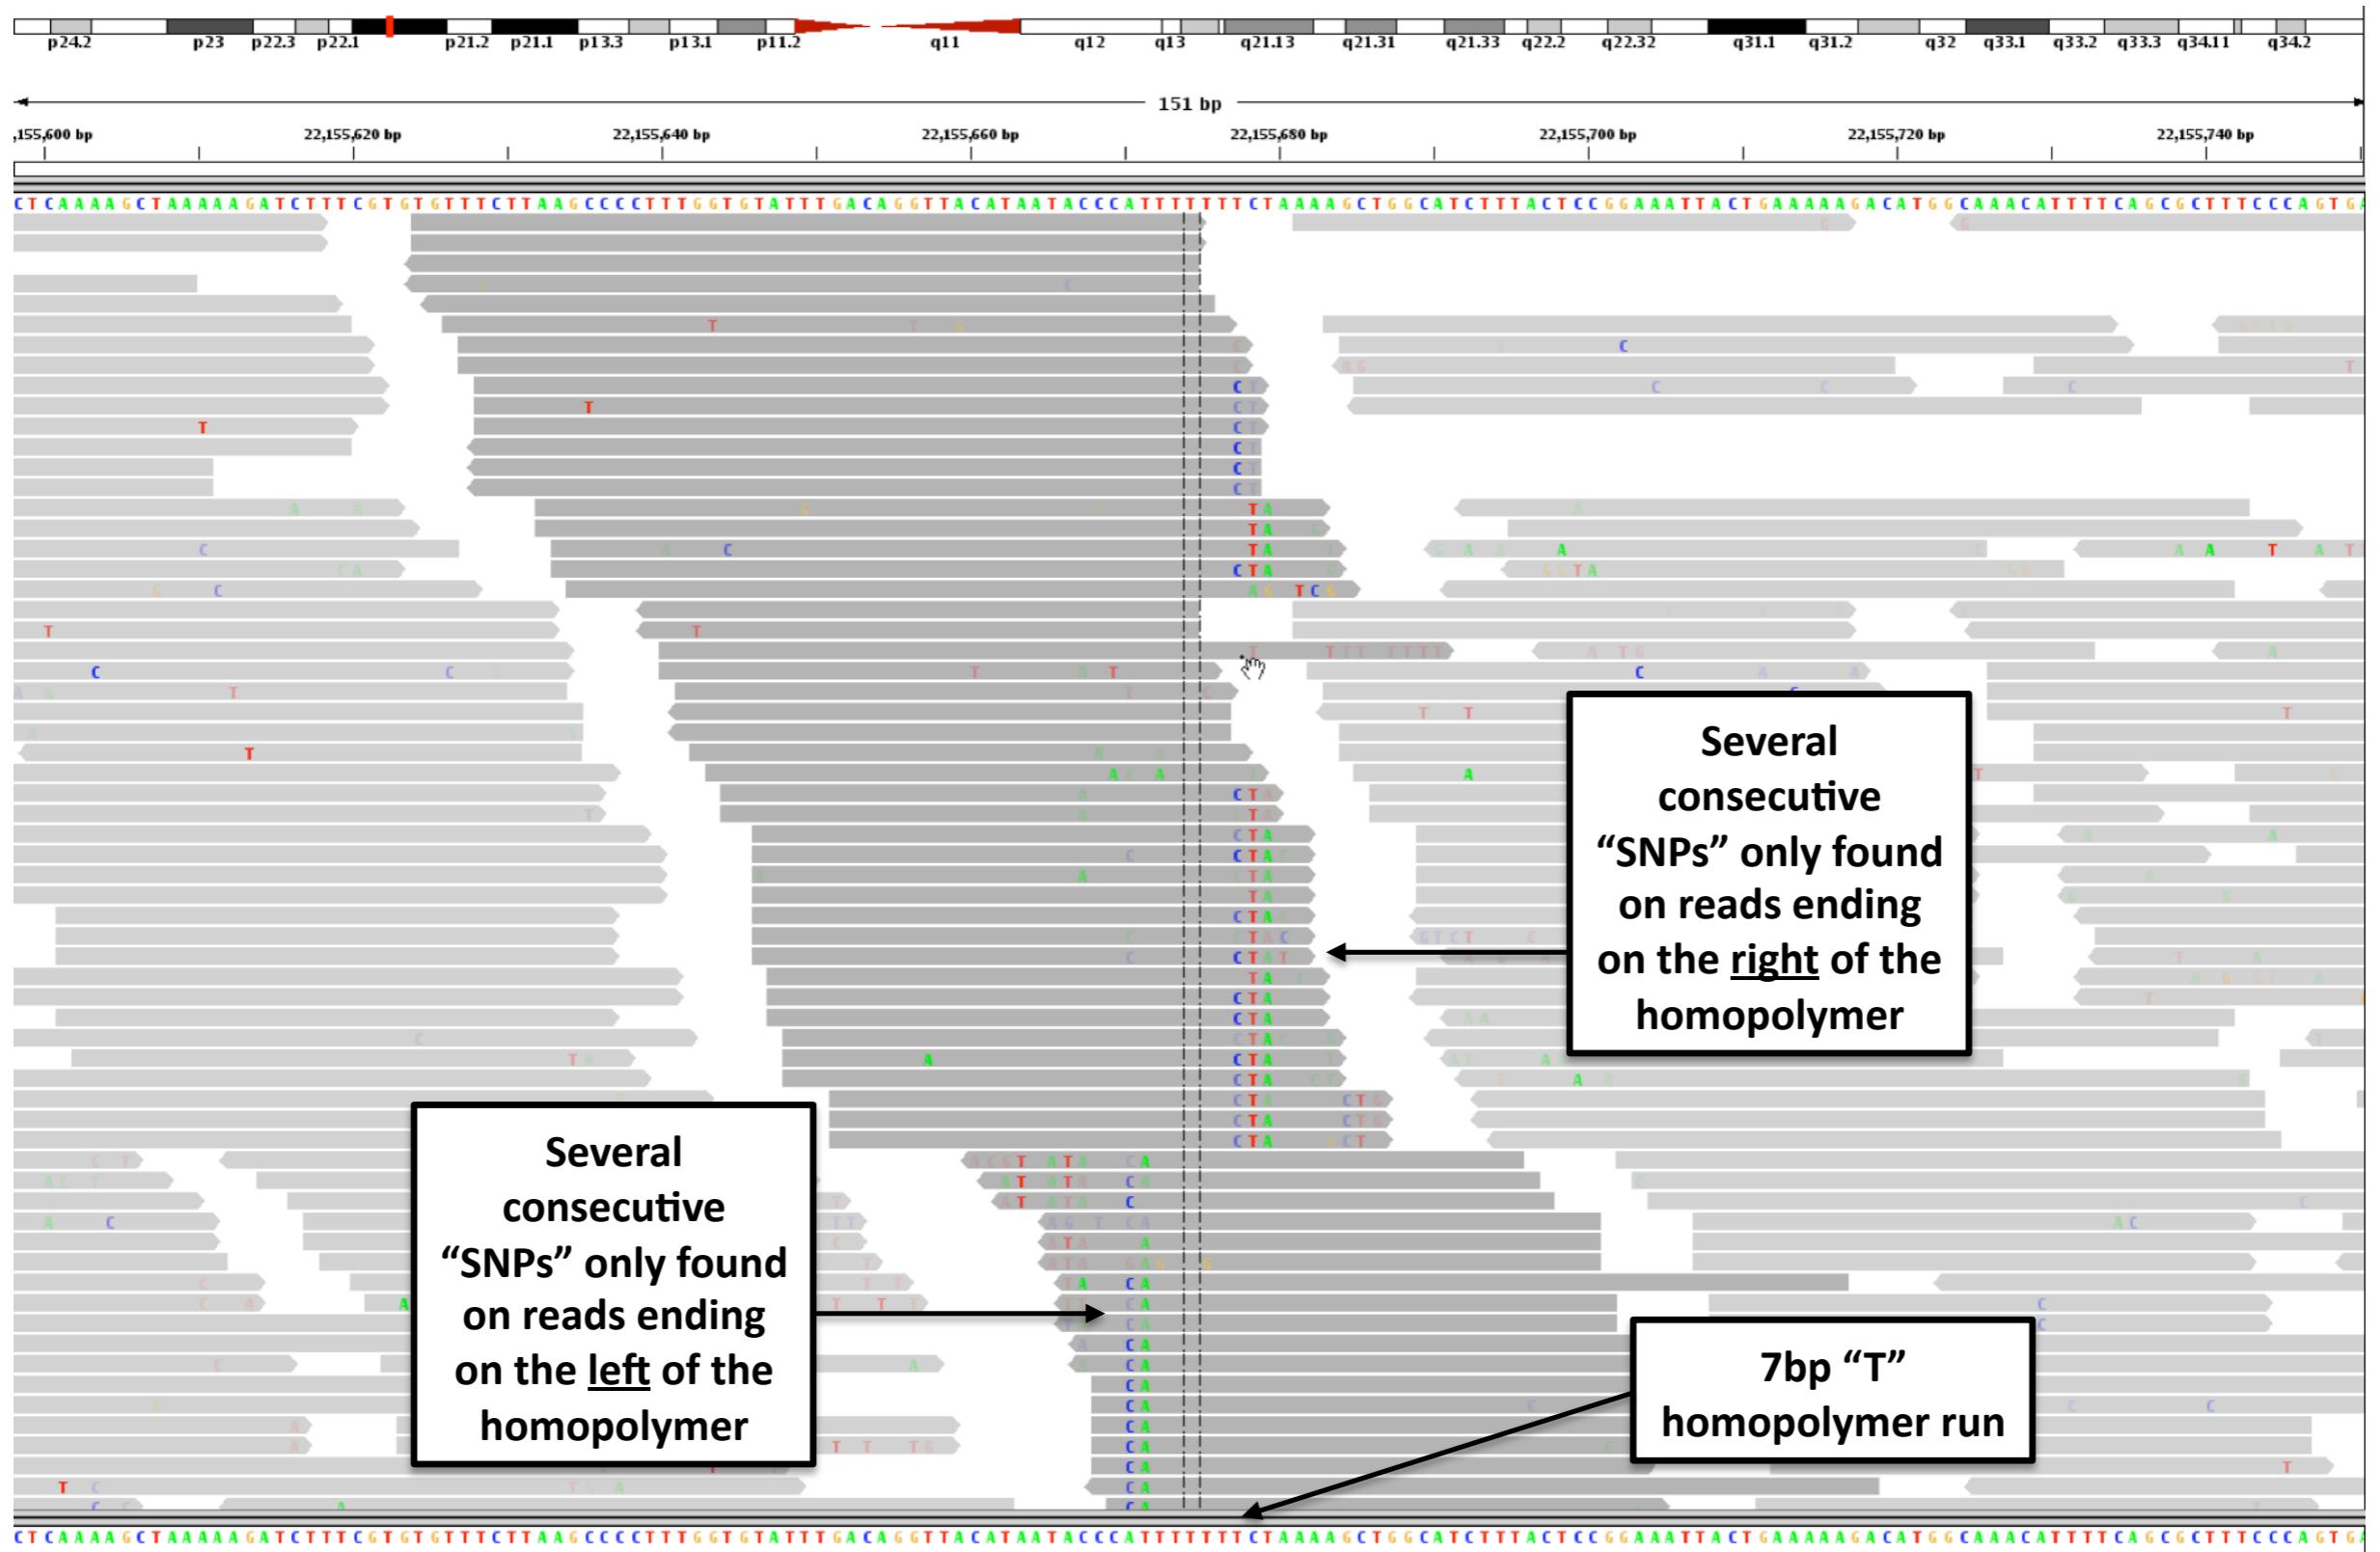


Why realign around indels?

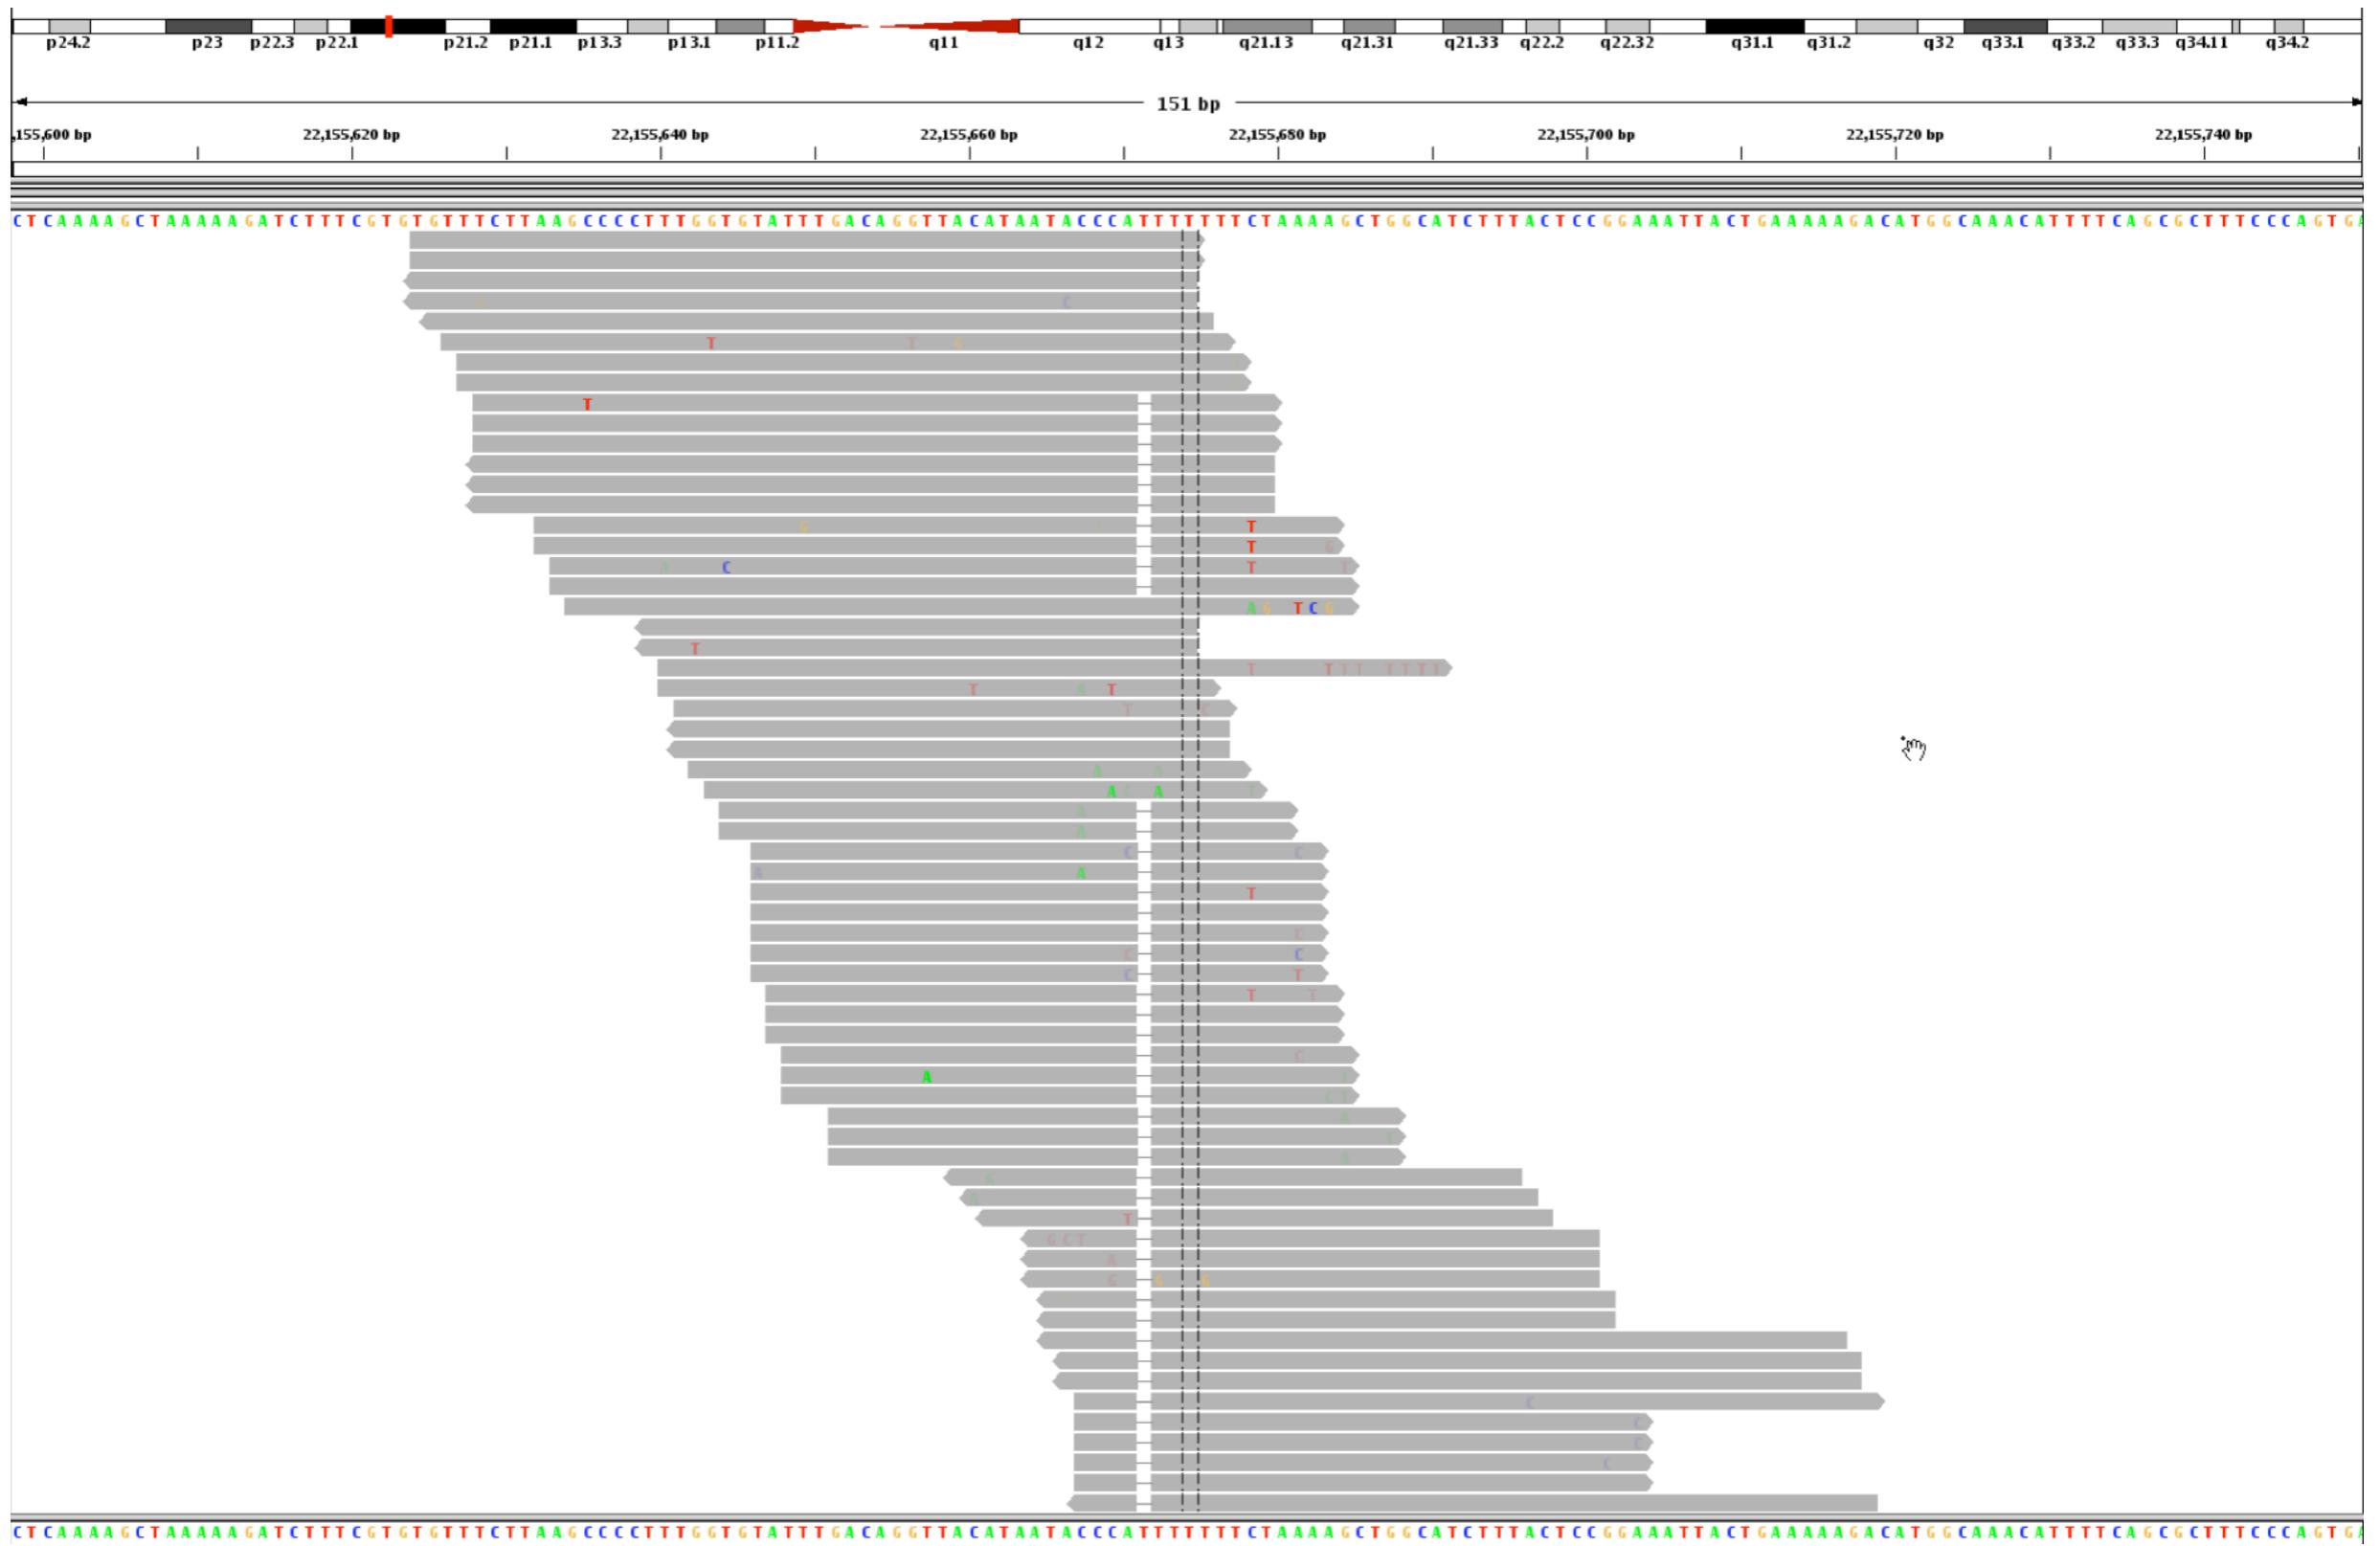
- InDels in reads (especially near the ends) can trick the mappers into mis-aligning with mismatches
- These artifactual mismatches can harm base quality recalibration and variant detection (unless a sophisticated caller like the Haplotype Caller is used)

Realignment around indels helps improve the accuracy of several of the downstream processing steps.

An example of a strand-discordant locus



Local realignment uncovers the hidden indel in these reads and eliminates all the potential FP SNPs

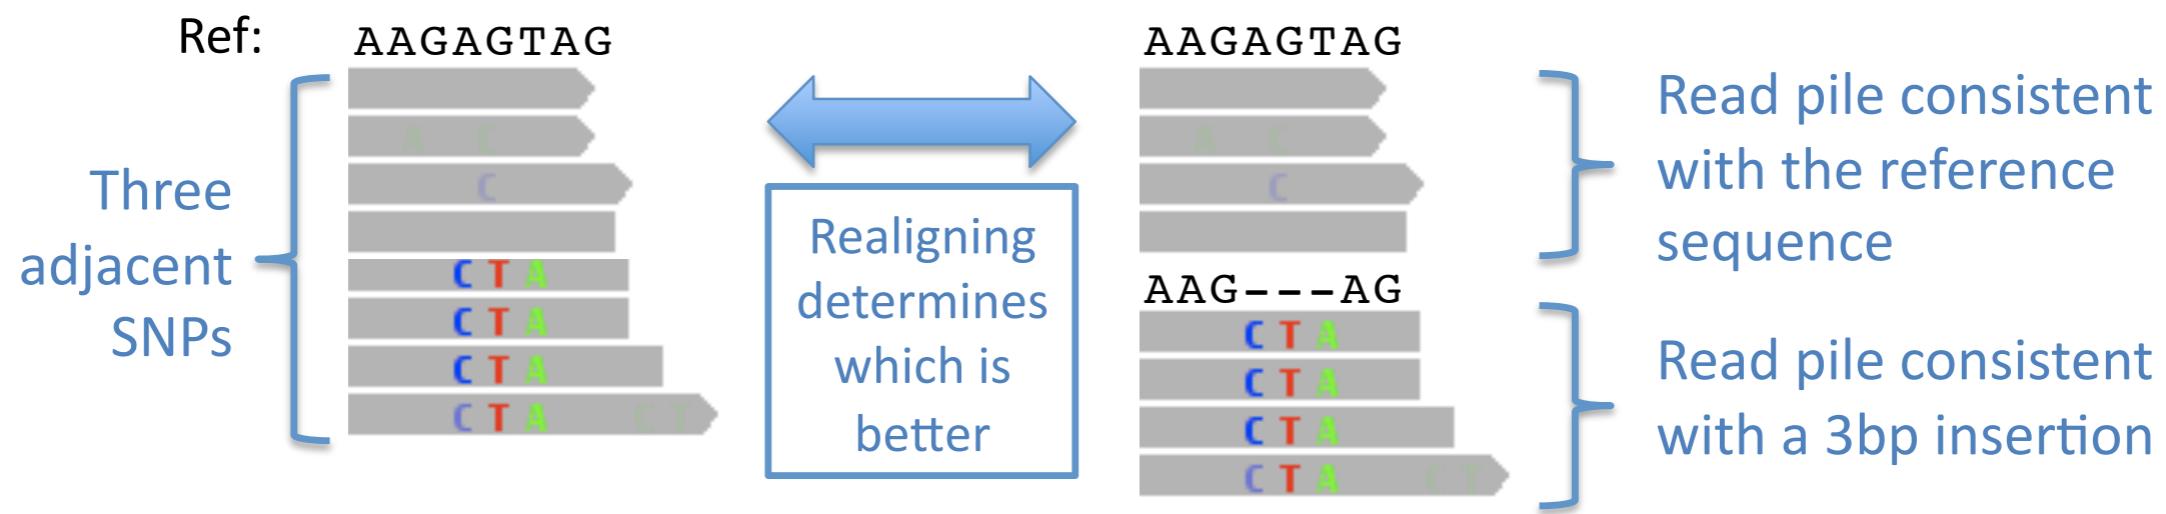


Three types of realignment targets

- Known sites (e.g. dbSNP, 1000 Genomes)
- Indels seen in original alignments (in CIGARs)
- Sites where evidence suggests a hidden indel
 - This is done from an entropy calculation
 - Computes activity score and if it is above a threshold, the region is included as realignment interval

Local realignment identifies most parsimonious alignment along all reads at a problematic locus

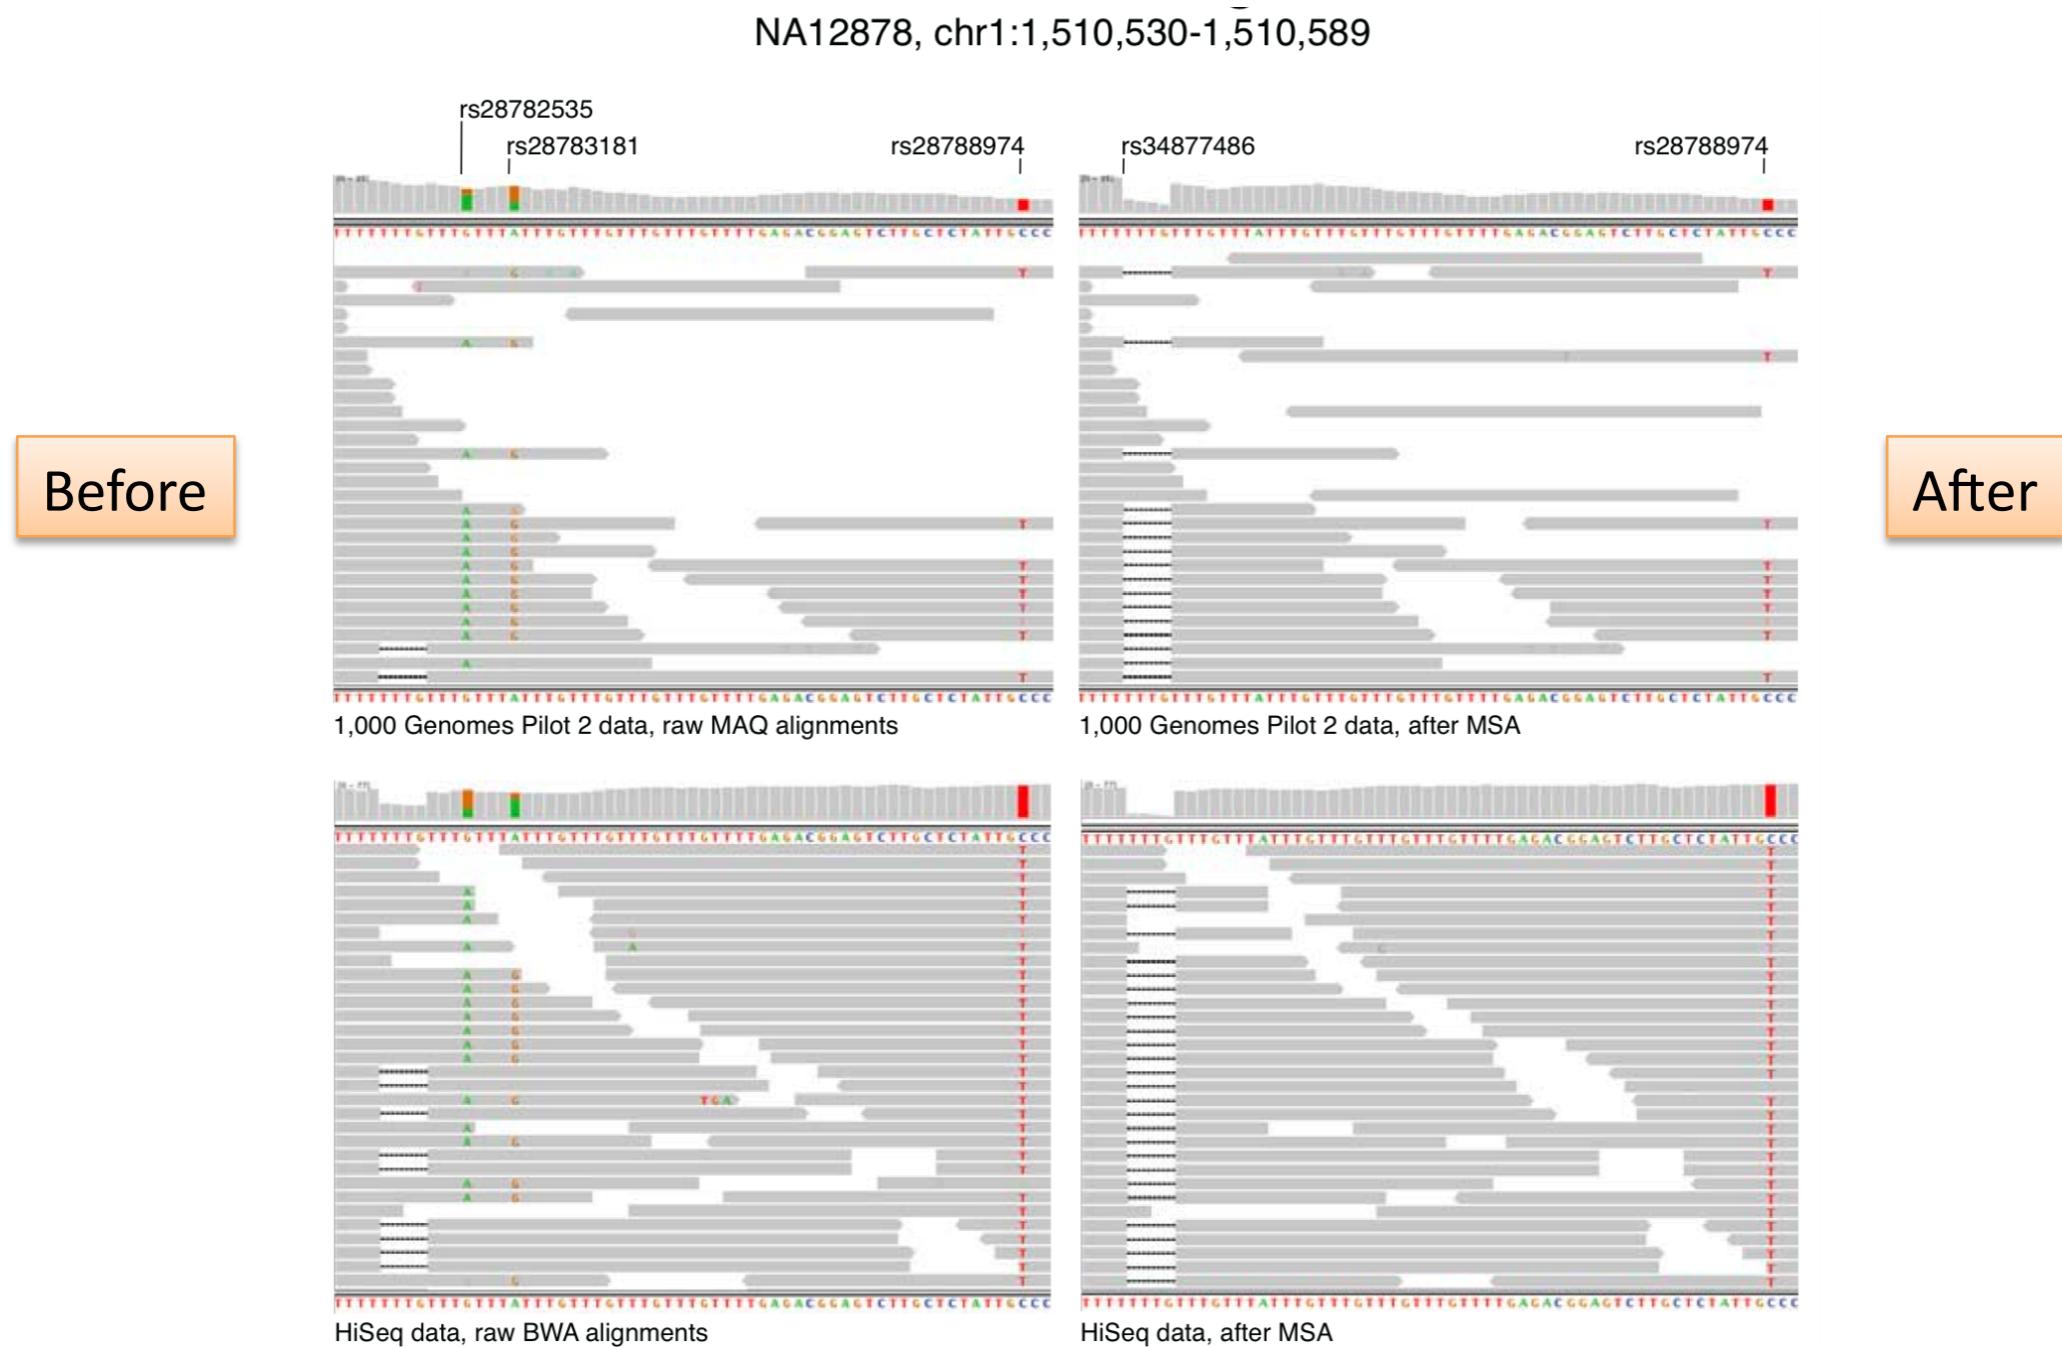
1. Find the best alternate consensus sequence that, together with the reference, best fits the reads in a pile (maximum of 1 indel)



2. The score for an alternate consensus is the total sum of the quality scores of mismatching bases

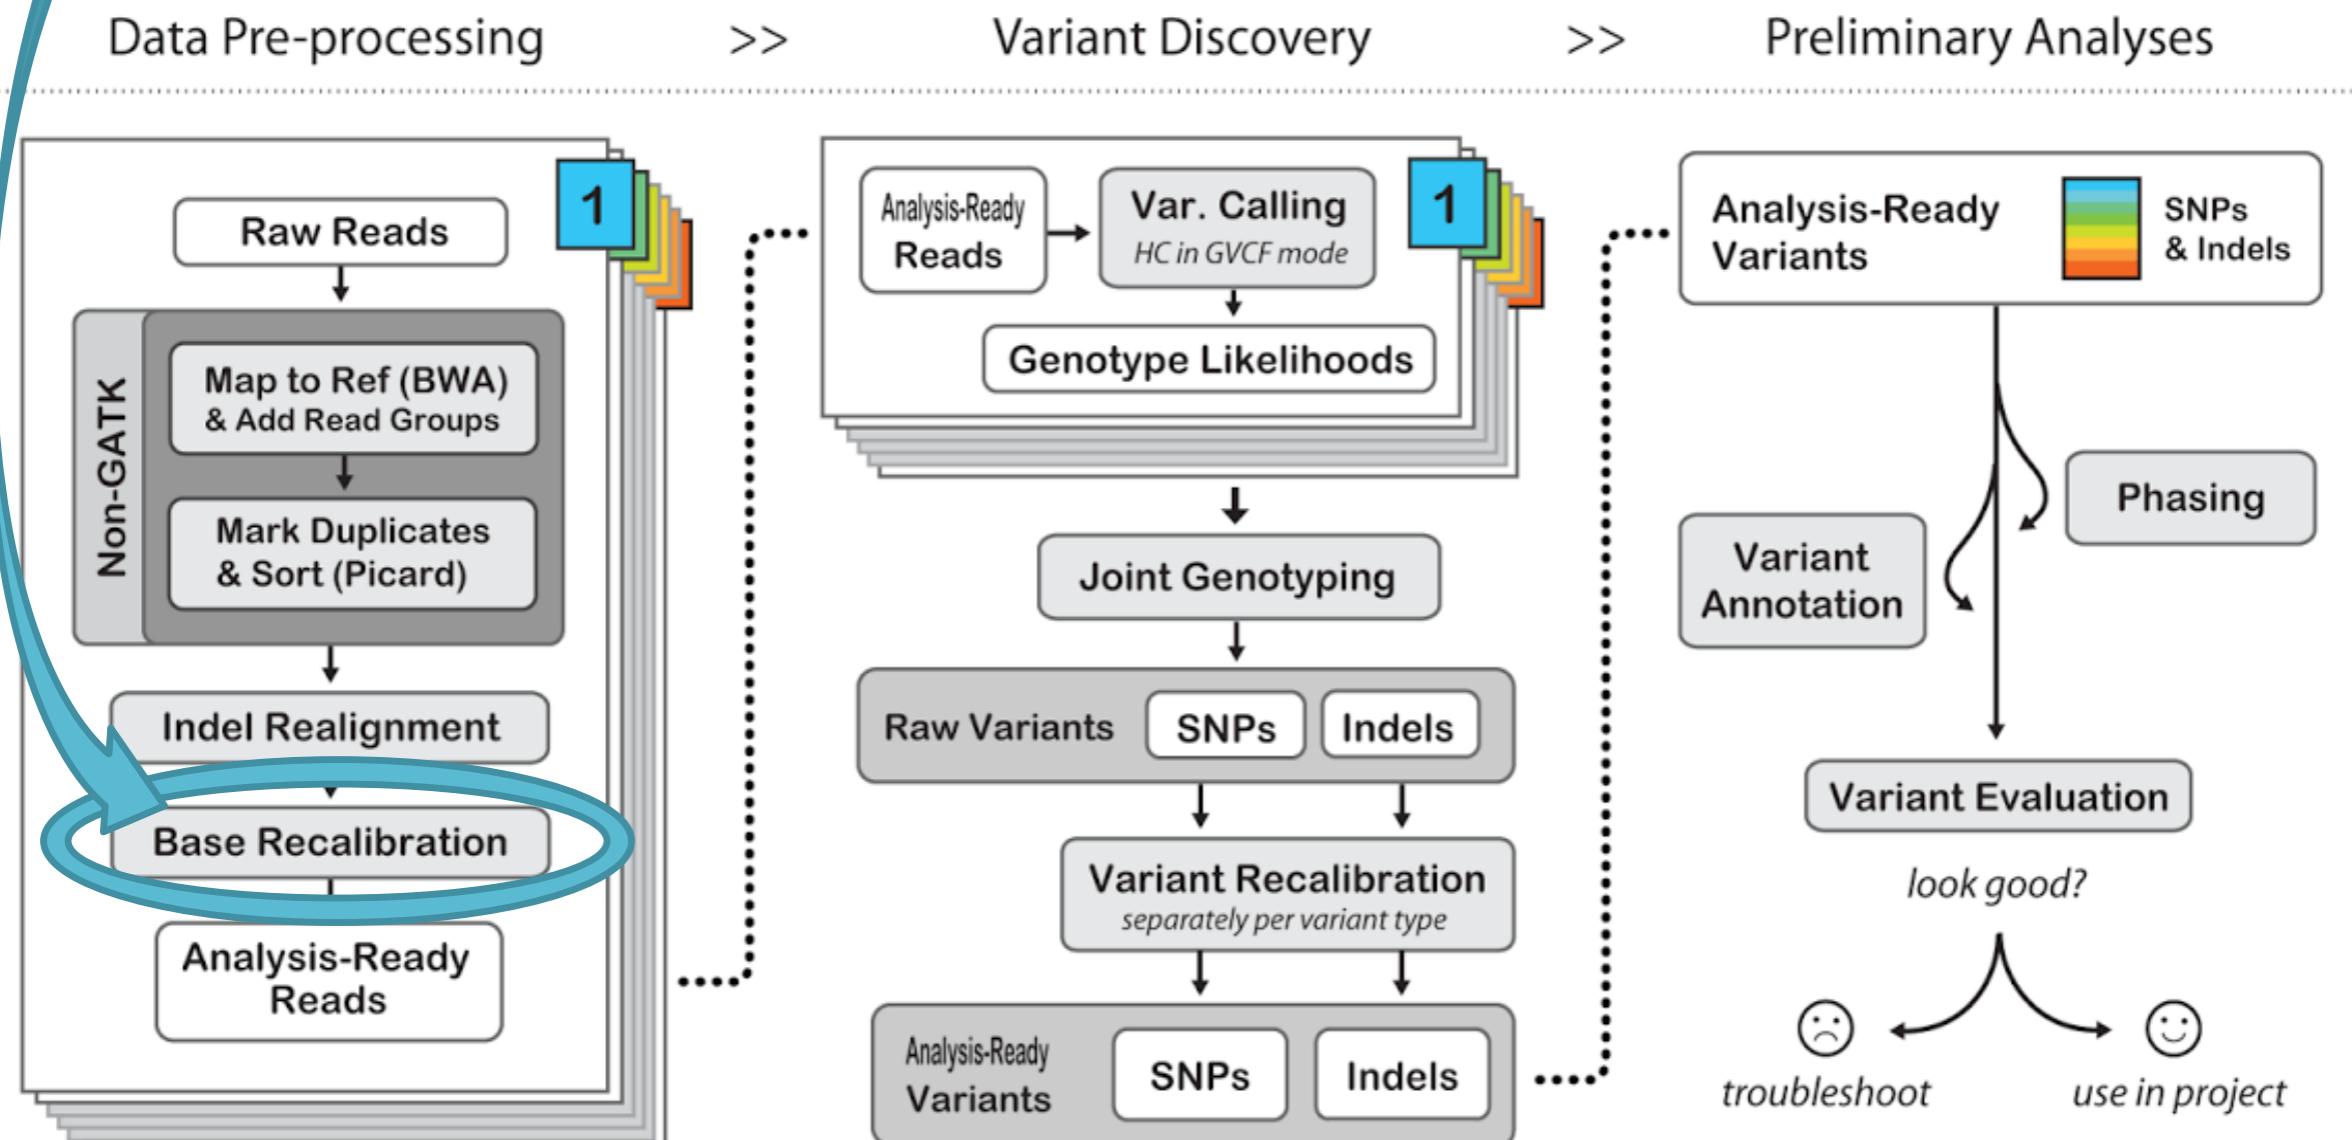
3. If the score of the best alternate consensus is sufficiently better than the original alignments (using a LOD score), then we accept the proposed realignment of the reads

This is what a realigned BAM looks like



We are here in the Best Practices workflow

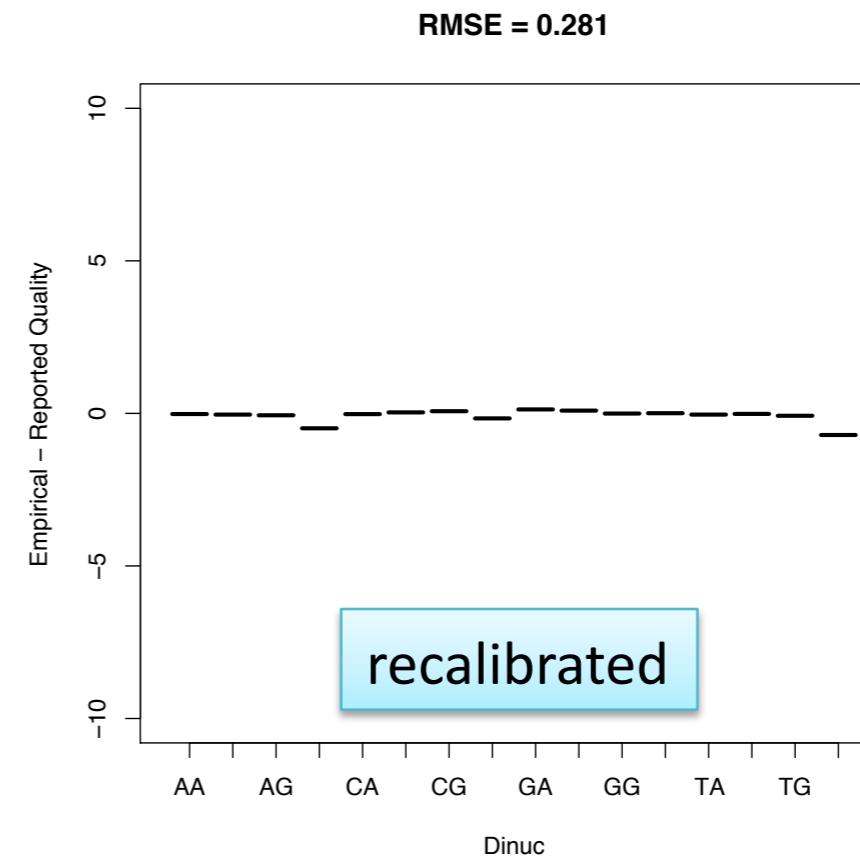
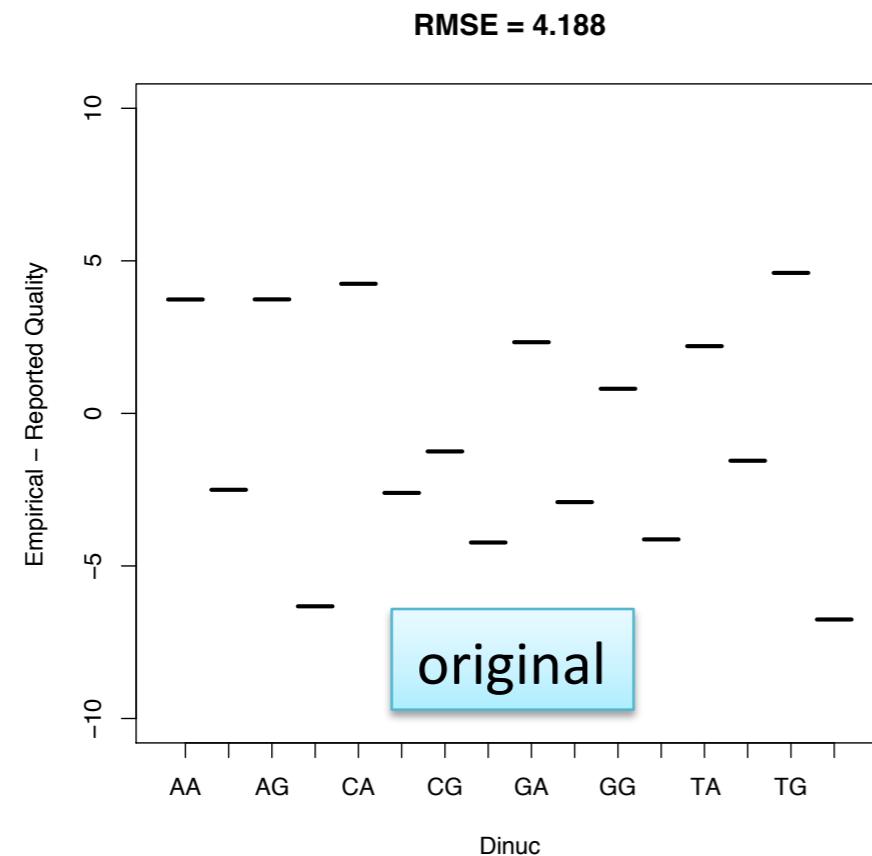
Base Recalibration



Quality scores issued by sequencers are **inaccurate** and **biased**

- Quality scores are critical for all downstream analysis
- Systematic biases are a major contributor to bad calls

Example of bias: qualities reported depending on nucleotide context

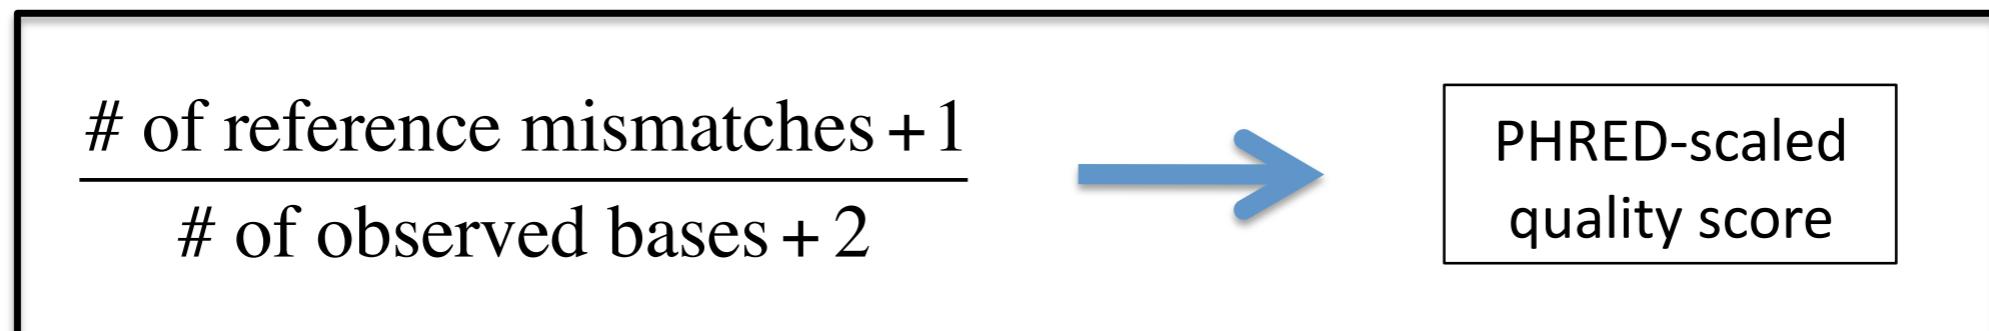


BQSR identifies patterns in how errors correlate with base features

- Empowered by looking at entire lane of data
- Analyze covariation among several features of a base, e.g.:
 - Reported quality score
 - Position within the read (machine cycle)
 - Preceding and current nucleotide (sequencing chemistry effect)
- Based on the patterns identified:
Apply corrections to recalibrate the quality scores of all reads in the BAM file.

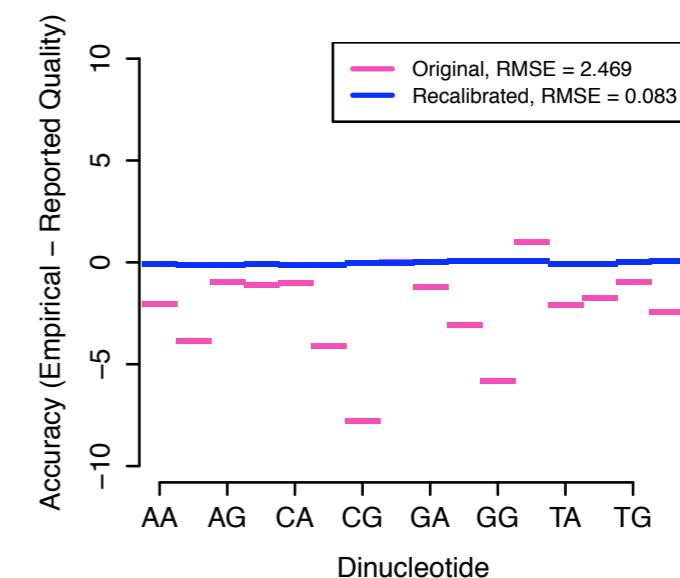
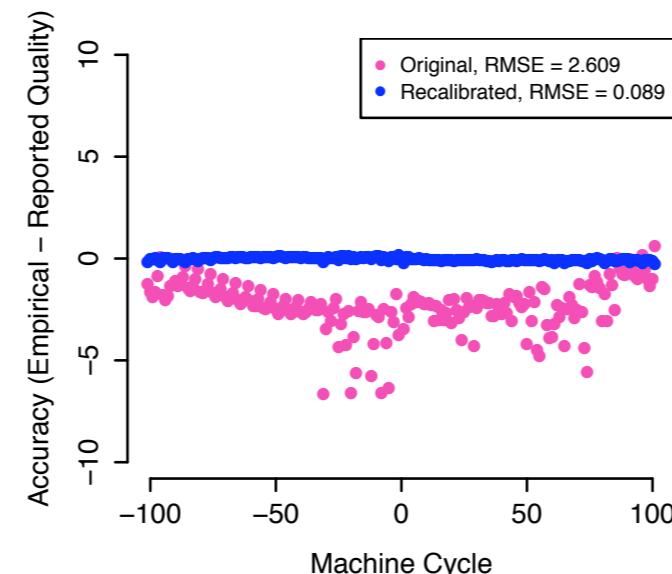
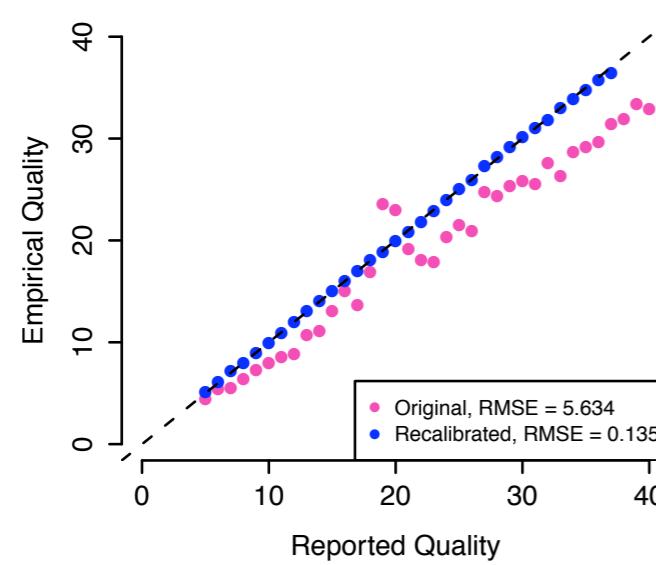
How covariates are analyzed to identify patterns

- Any sequence mismatch = error *except known variants!*
- Keep track of number of observations and number of errors as a function of various error covariates
(lane, original quality score, machine cycle, and sequencing context)

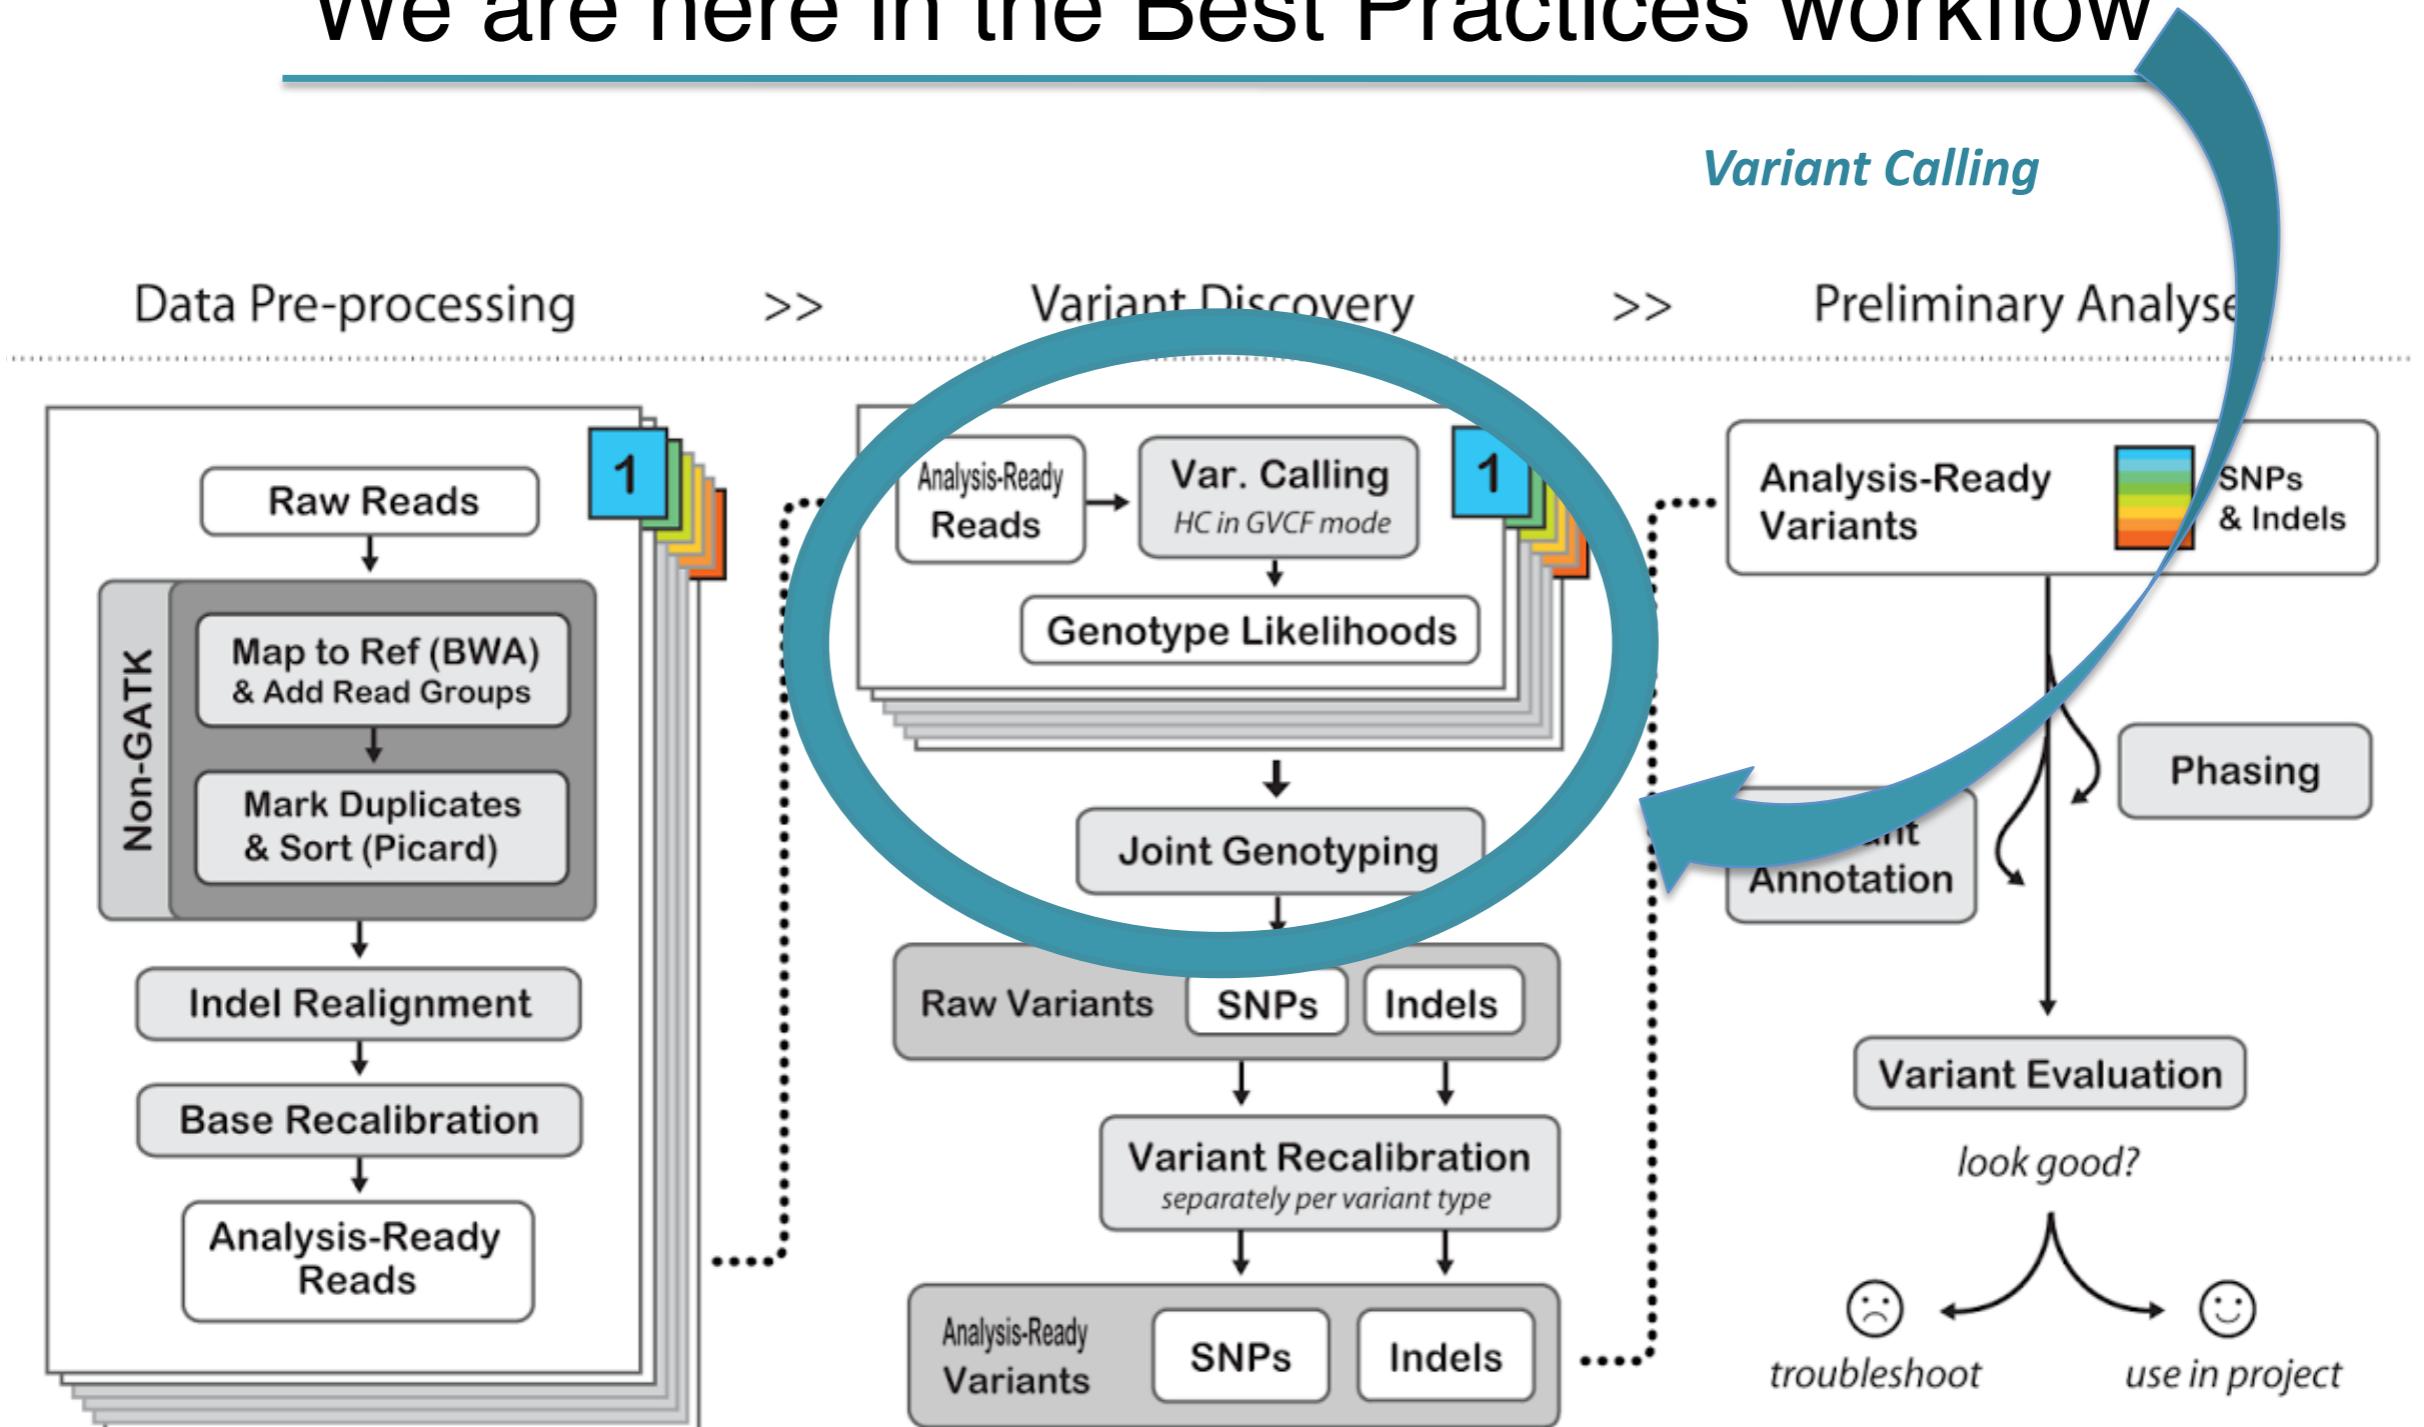


Did the recalibration work properly?

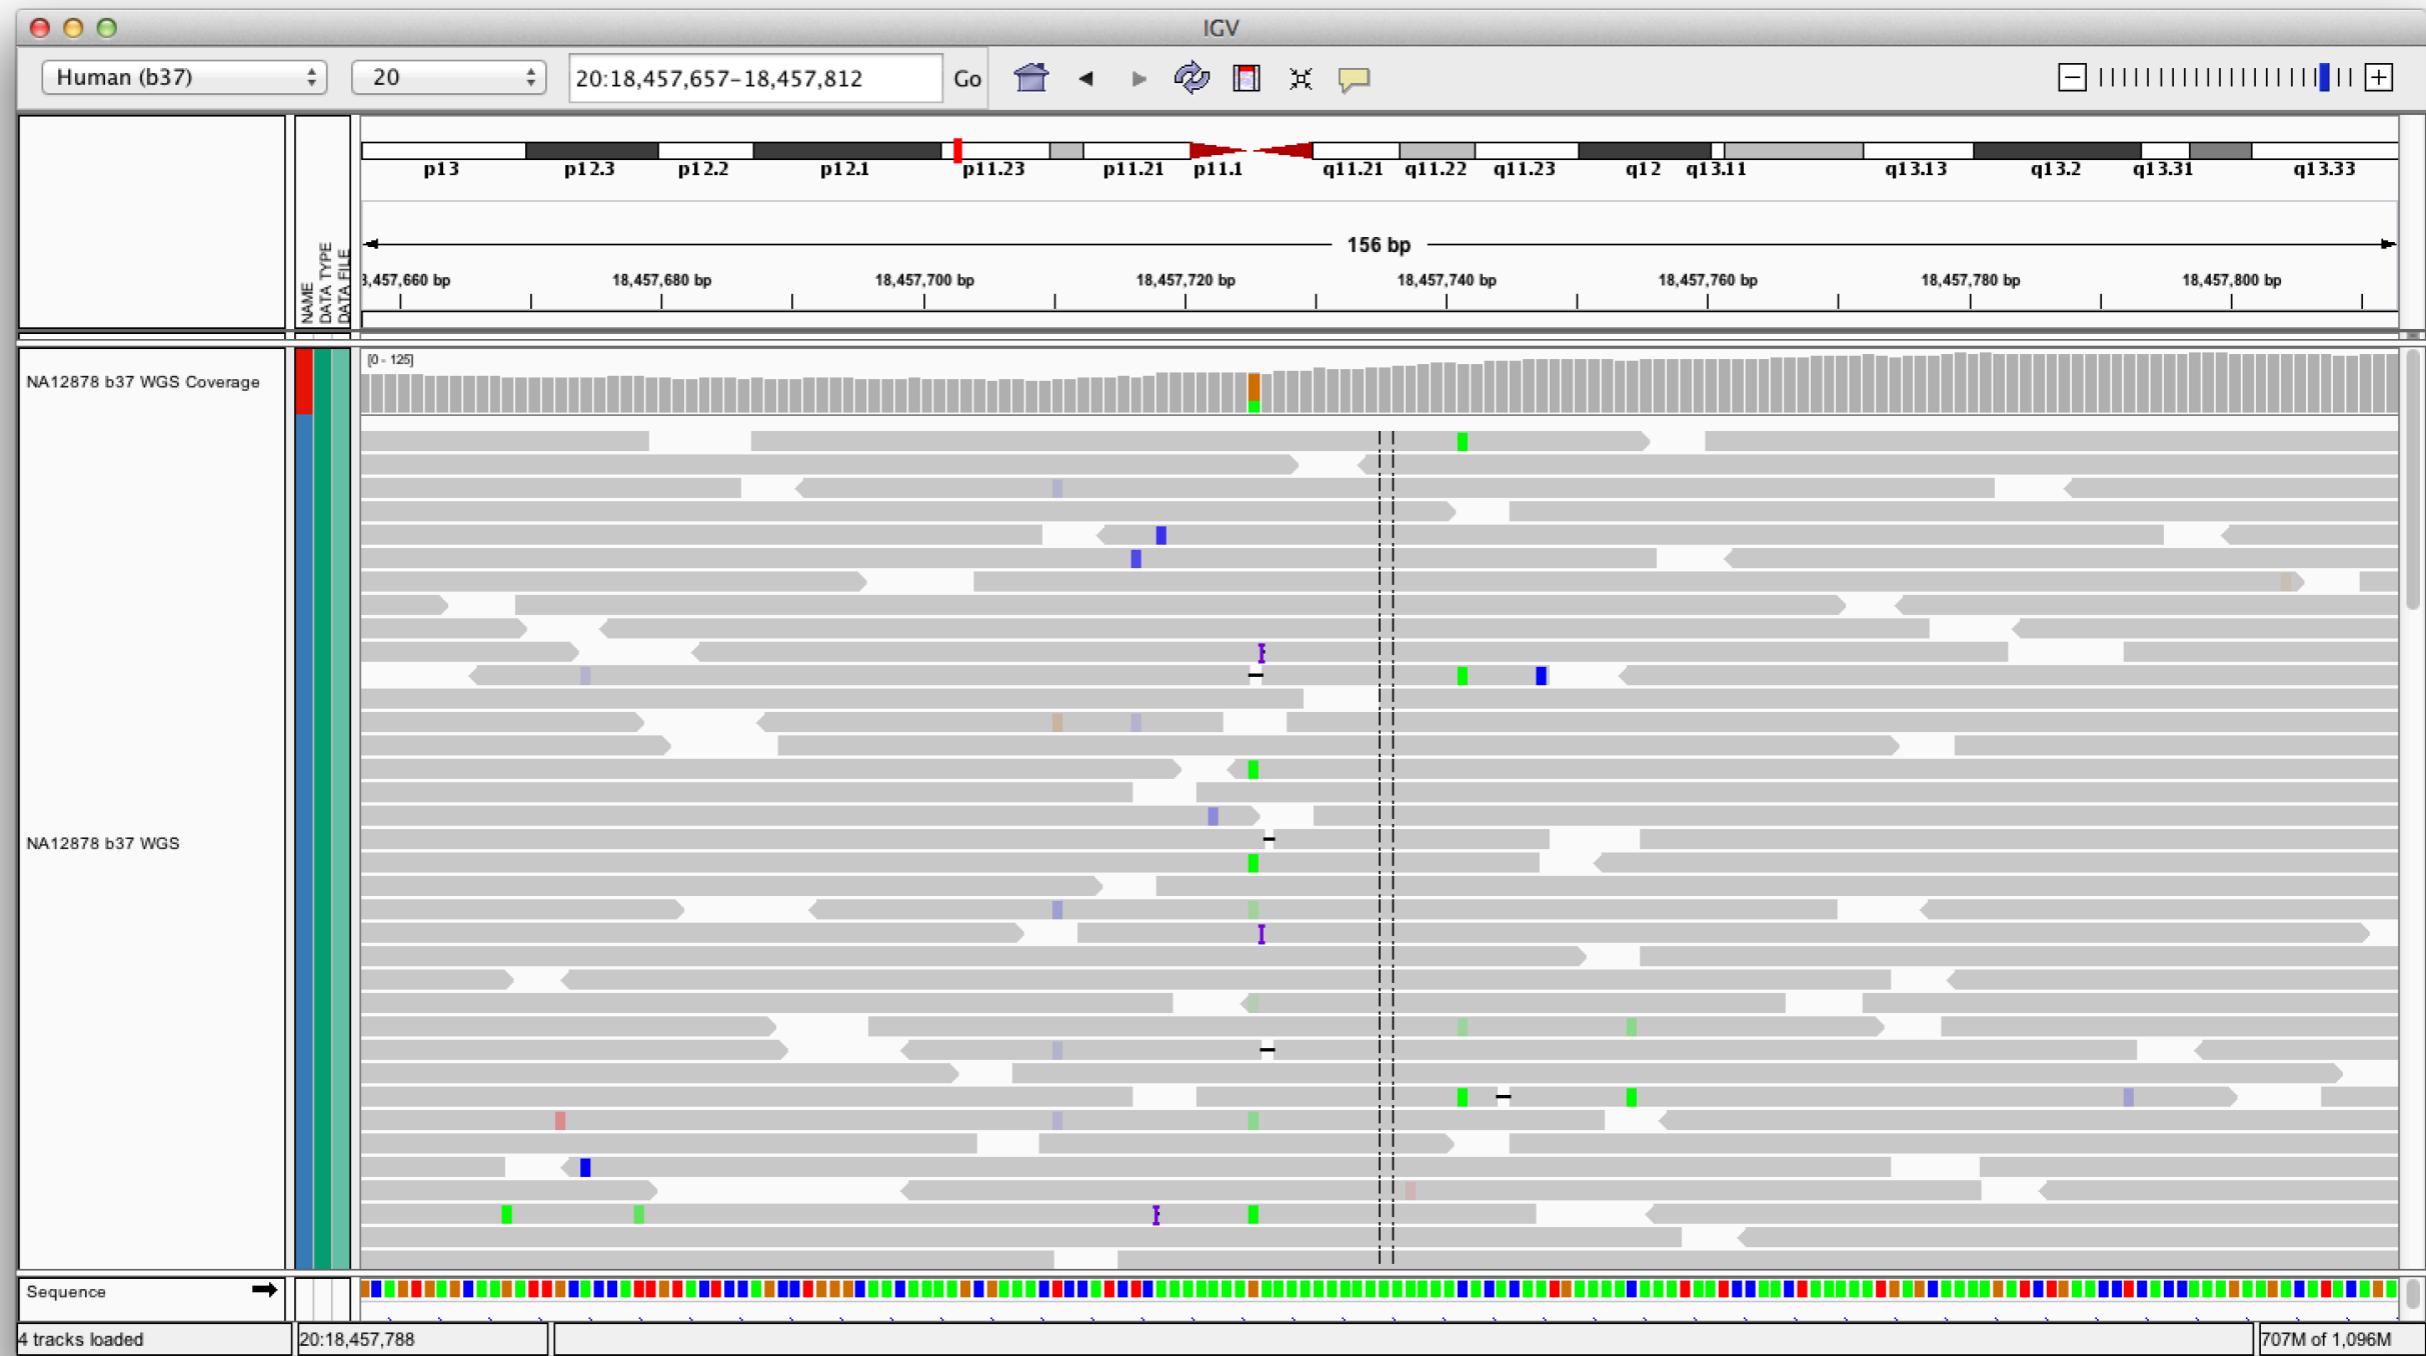
Post-recalibration quality scores should fit the empirically-derived quality scores very well; no obvious systematic biases should remain



We are here in the Best Practices workflow



Real mutations are hidden in the noise



Summed up in GATK terms

Bayesian model

$$\Pr\{G|D\} = \frac{\Pr\{G\} \Pr\{D|G\}}{\sum_i \Pr\{G_i\} \Pr\{D|G_i\}}, \text{ [Bayes' rule]}$$
$$\Pr\{D|G\} = \prod_j \left(\frac{\Pr\{D_j|H_1\}}{2} + \frac{\Pr\{D_j|H_2\}}{2} \right) \text{ where } G = H_1H_2$$

$\Pr\{D|H\}$ is the haploid likelihood function

Prior of the genotype Likelihood of the genotype

Diploid assumption

Variant callers in GATK

- **UnifiedGenotyper**

Call SNPs and indels separately by considering each variant locus independently

- Accepts any ploidy
- Pooled calling

- **HaplotypeCaller**

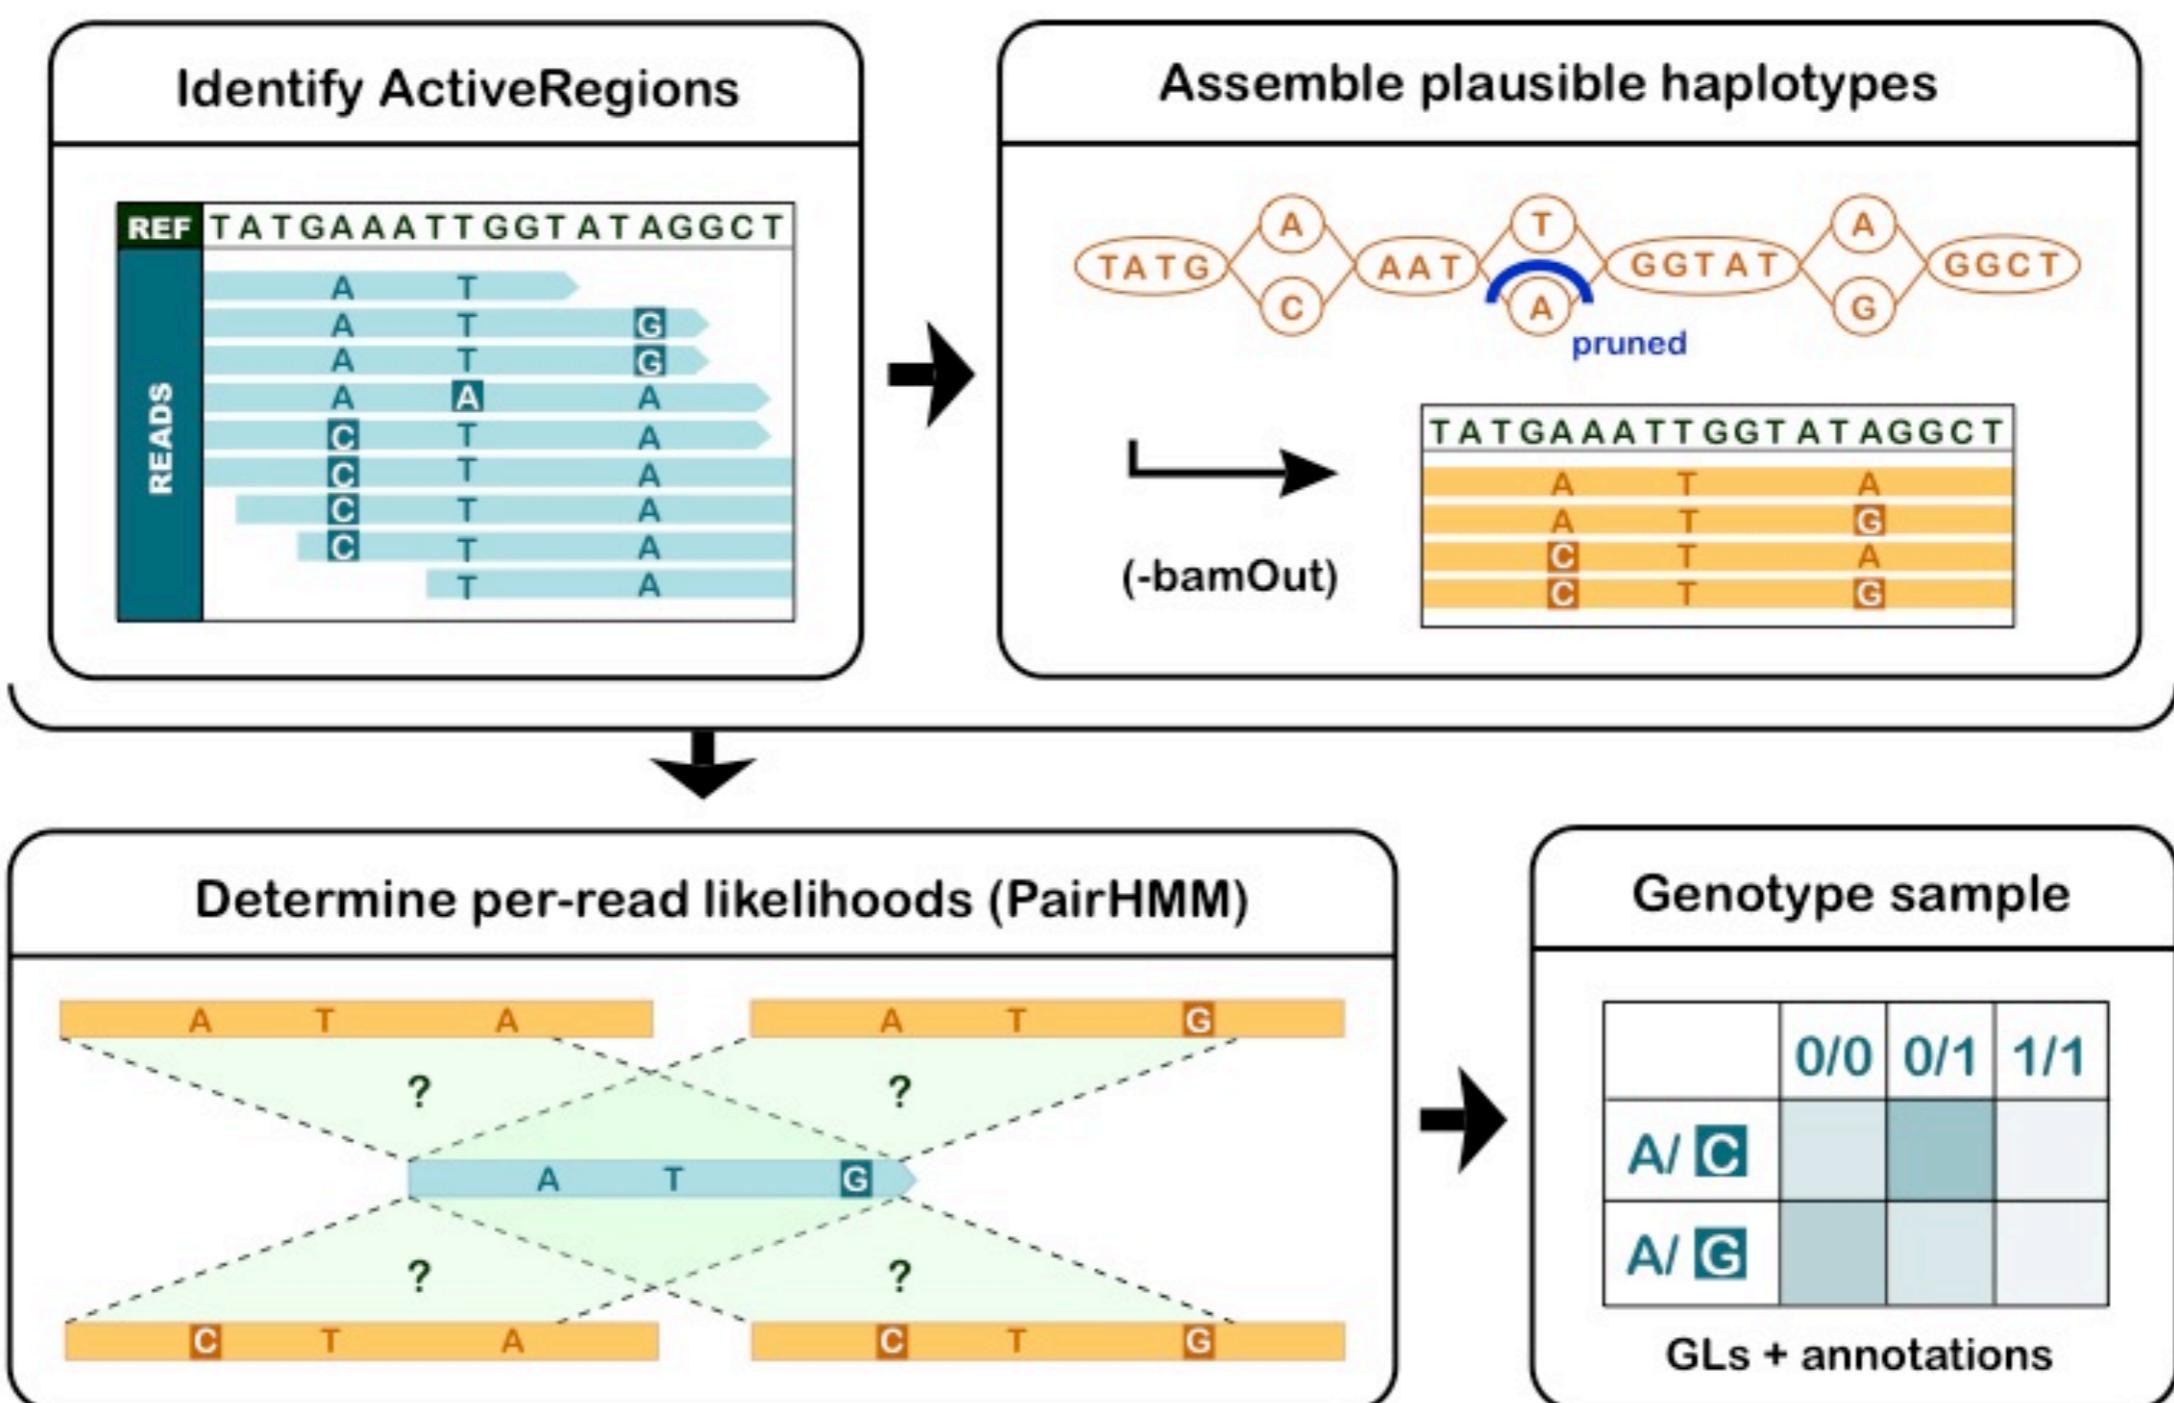
Call SNPs, indels, and some SVs simultaneously by doing local re-assembly and considering haplotypes

- More accurate (esp. indels)
- Reference confidence model
- Replaces UG

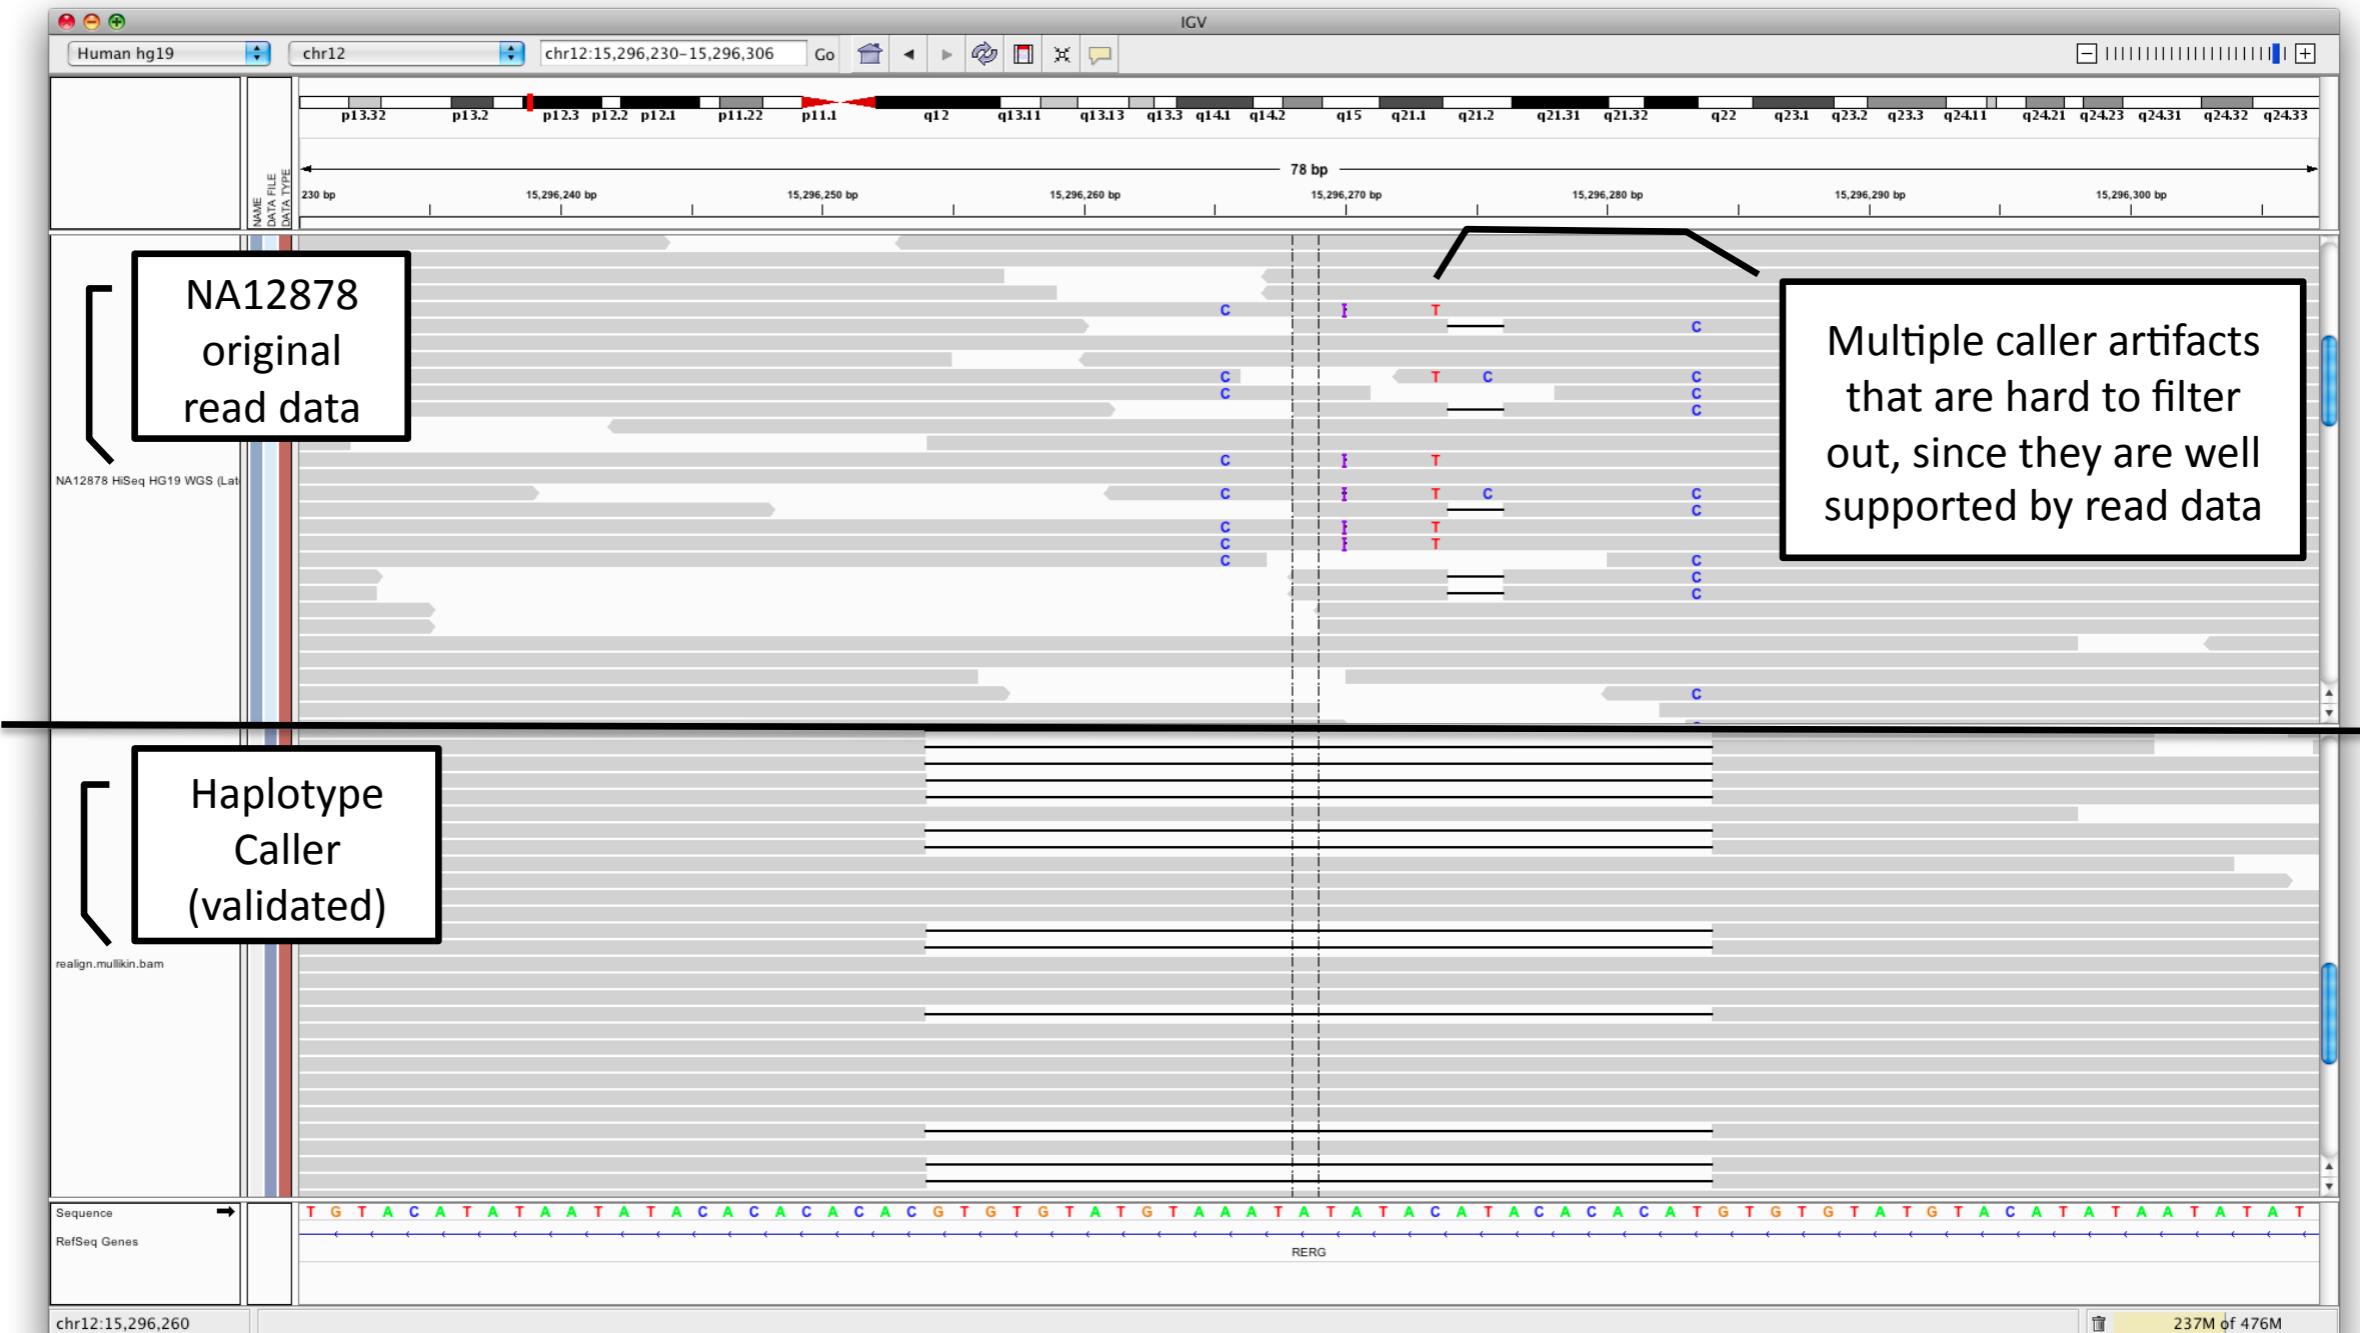
HaplotypeCaller method overview

- Call SNPs, indels, and some SVs simultaneously by doing local re-assembly and considering haplotypes
 - Determine if a region has **potential variation**
 - Make **deBruijn assembly graph** of the region
 - Paths in the graph = **potential haplotypes** to evaluate
 - Calculate **data likelihoods** given the haplotypes (PairHMM)
 - **Identify variants** on most likely haplotypes
 - Compute **allele frequency distribution** to determine most likely allele count, and emit a variant call if appropriate
 - If emitting a variant, **assign genotype** to each sample

HC method illustrated

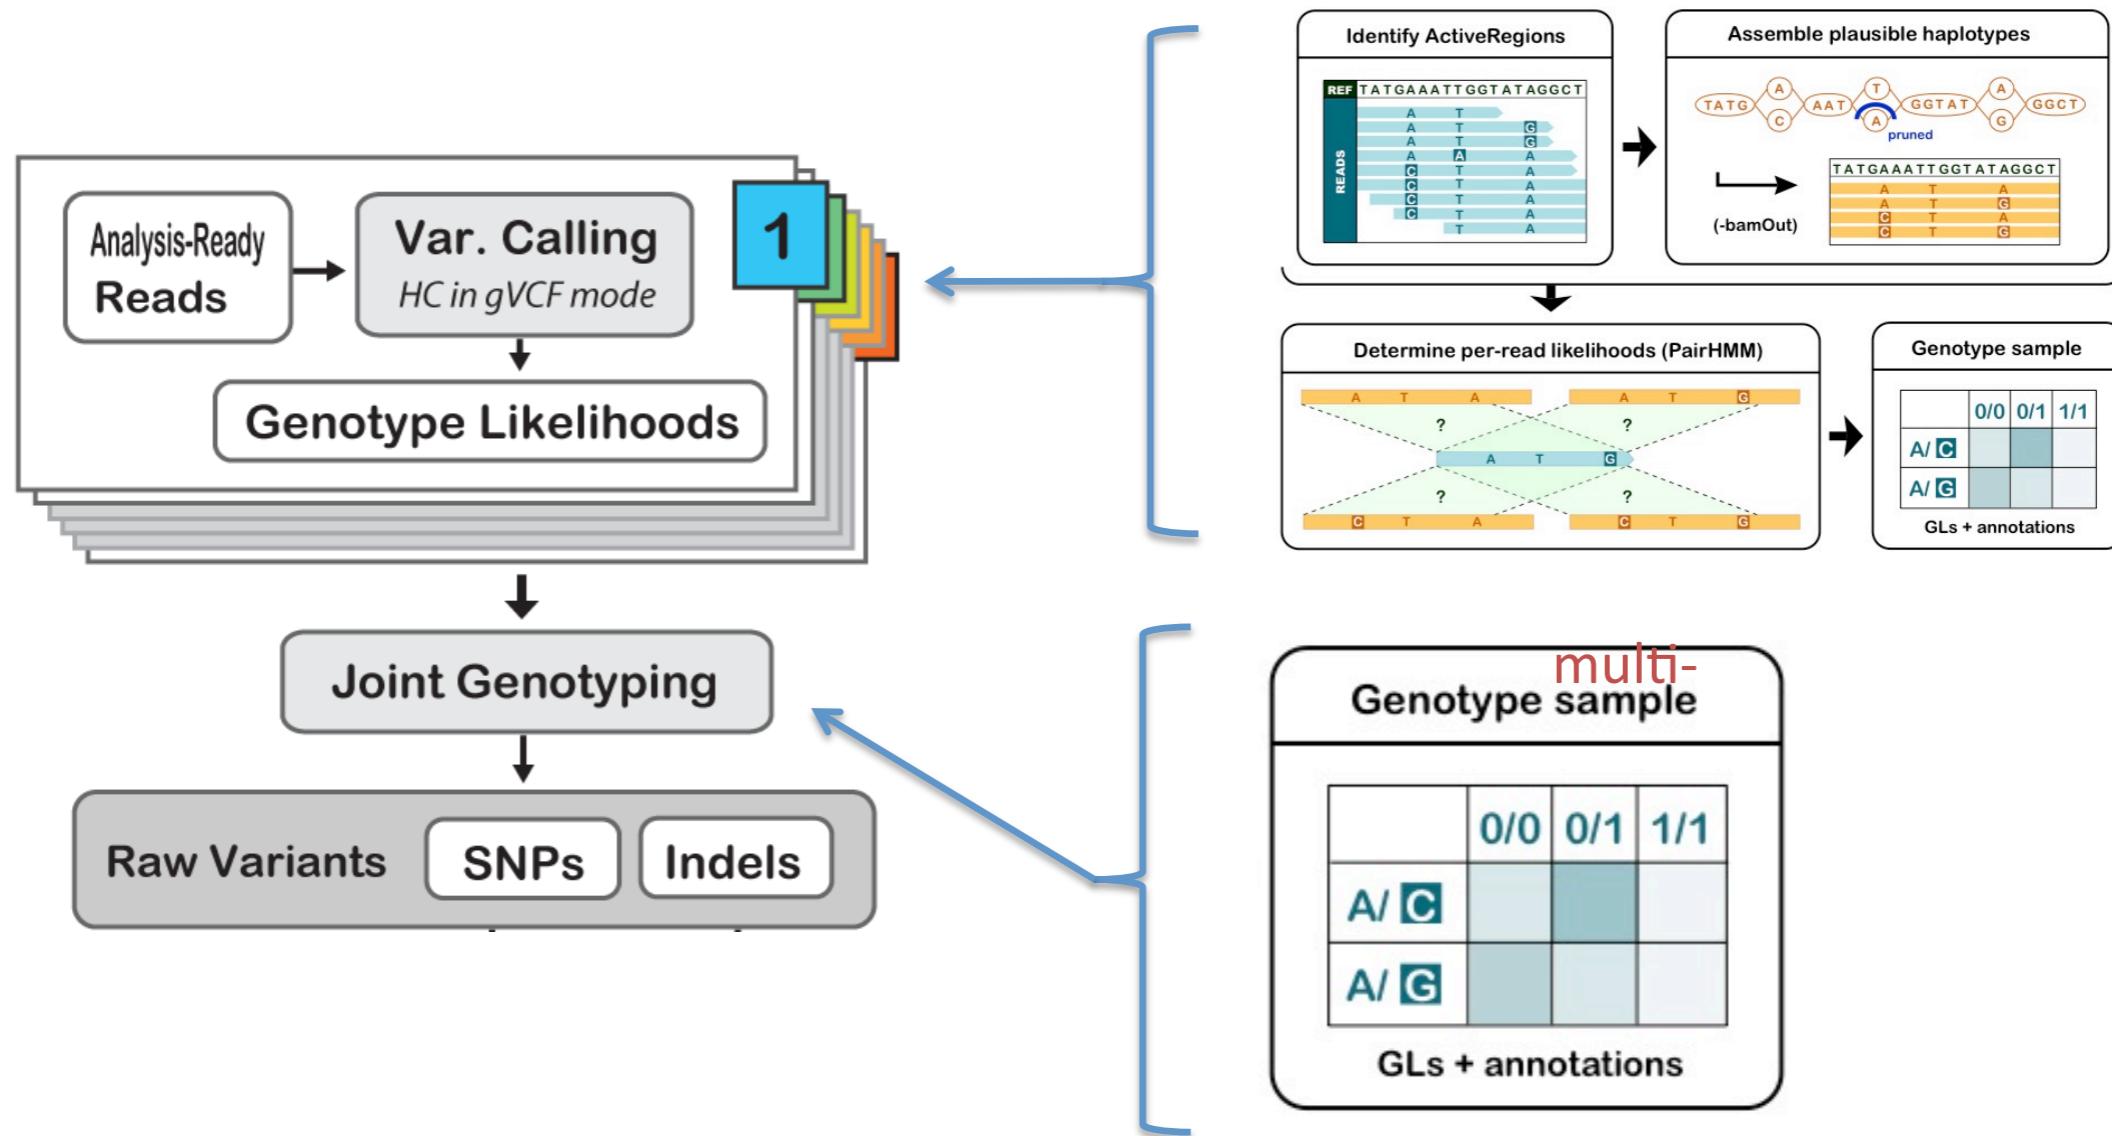


Artifactual SNPs and small indels caused by large indel can be recovered by assembly

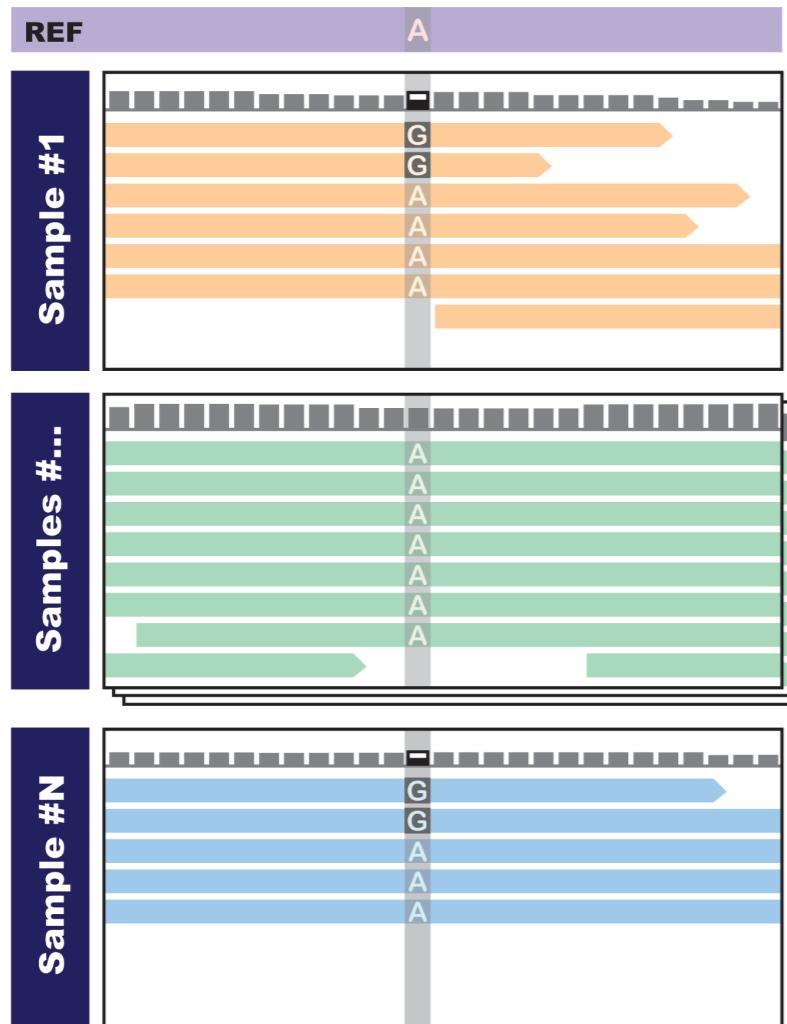


Joint Genotyping

Add a joint analysis step to take advantage
of cohort / pop genetics data



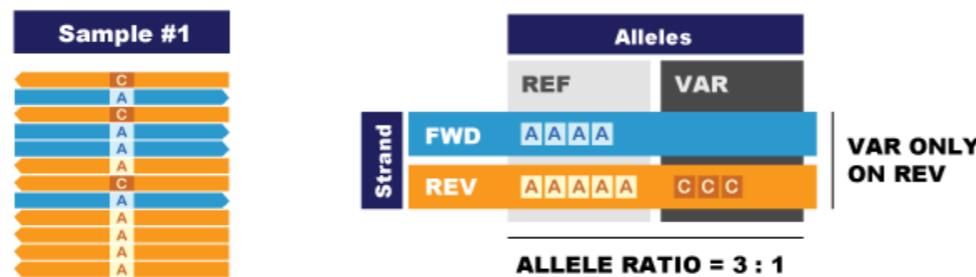
Joint discovery empowers discovery at difficult sites



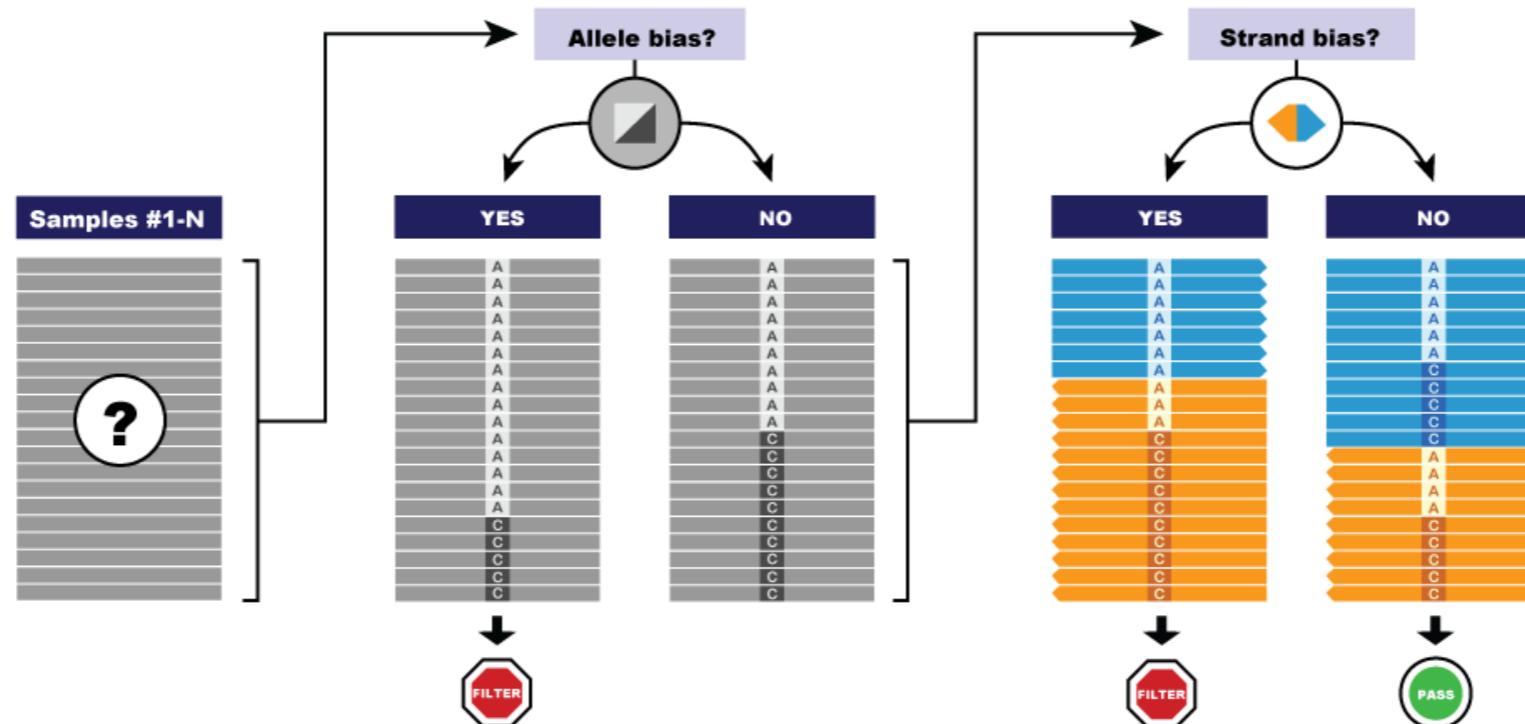
- If we analyze Sample #1 or Sample #N alone we are not confident that the variant is real
- If we see both samples then we are more confident that there is real variation at this site in the cohort

Joint discovery helps resolve bias issues

A. Single sample showing strand and allelic biases

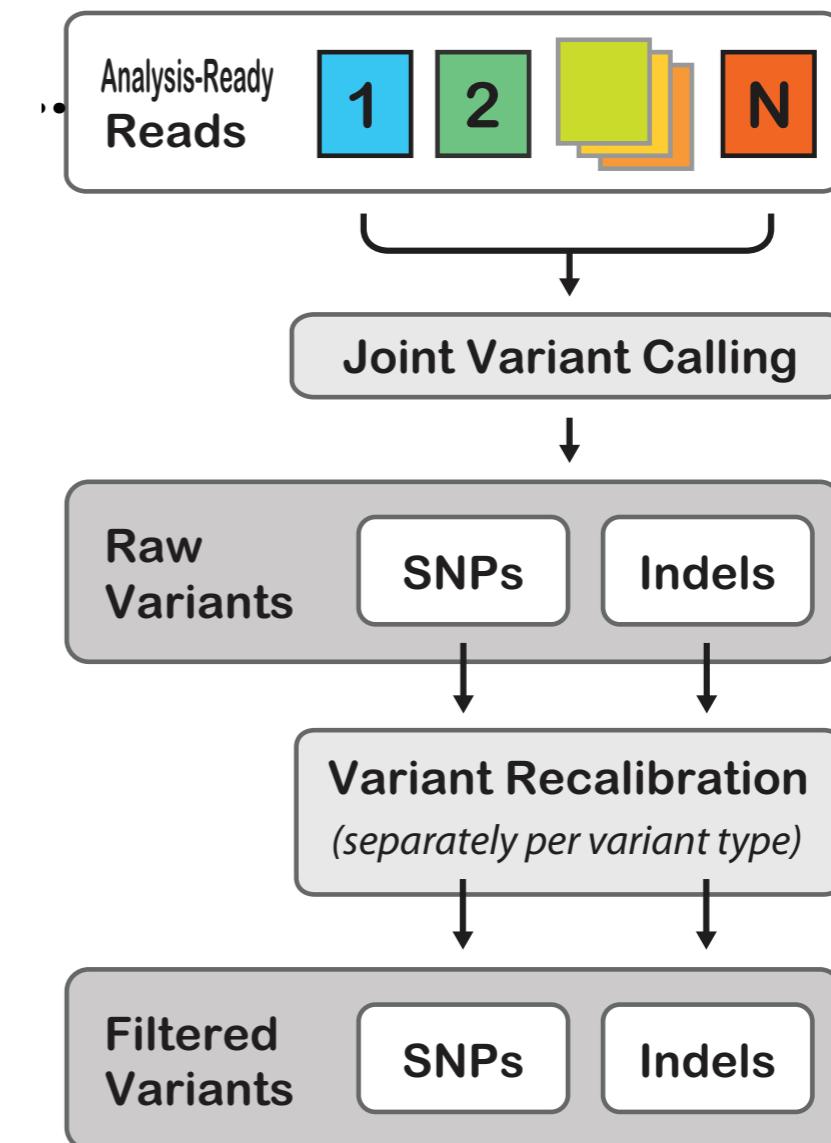


B. Decision process using evidence from multiple samples to filter out sites showing systematic biases



Classic approach to multi-sample variant discovery

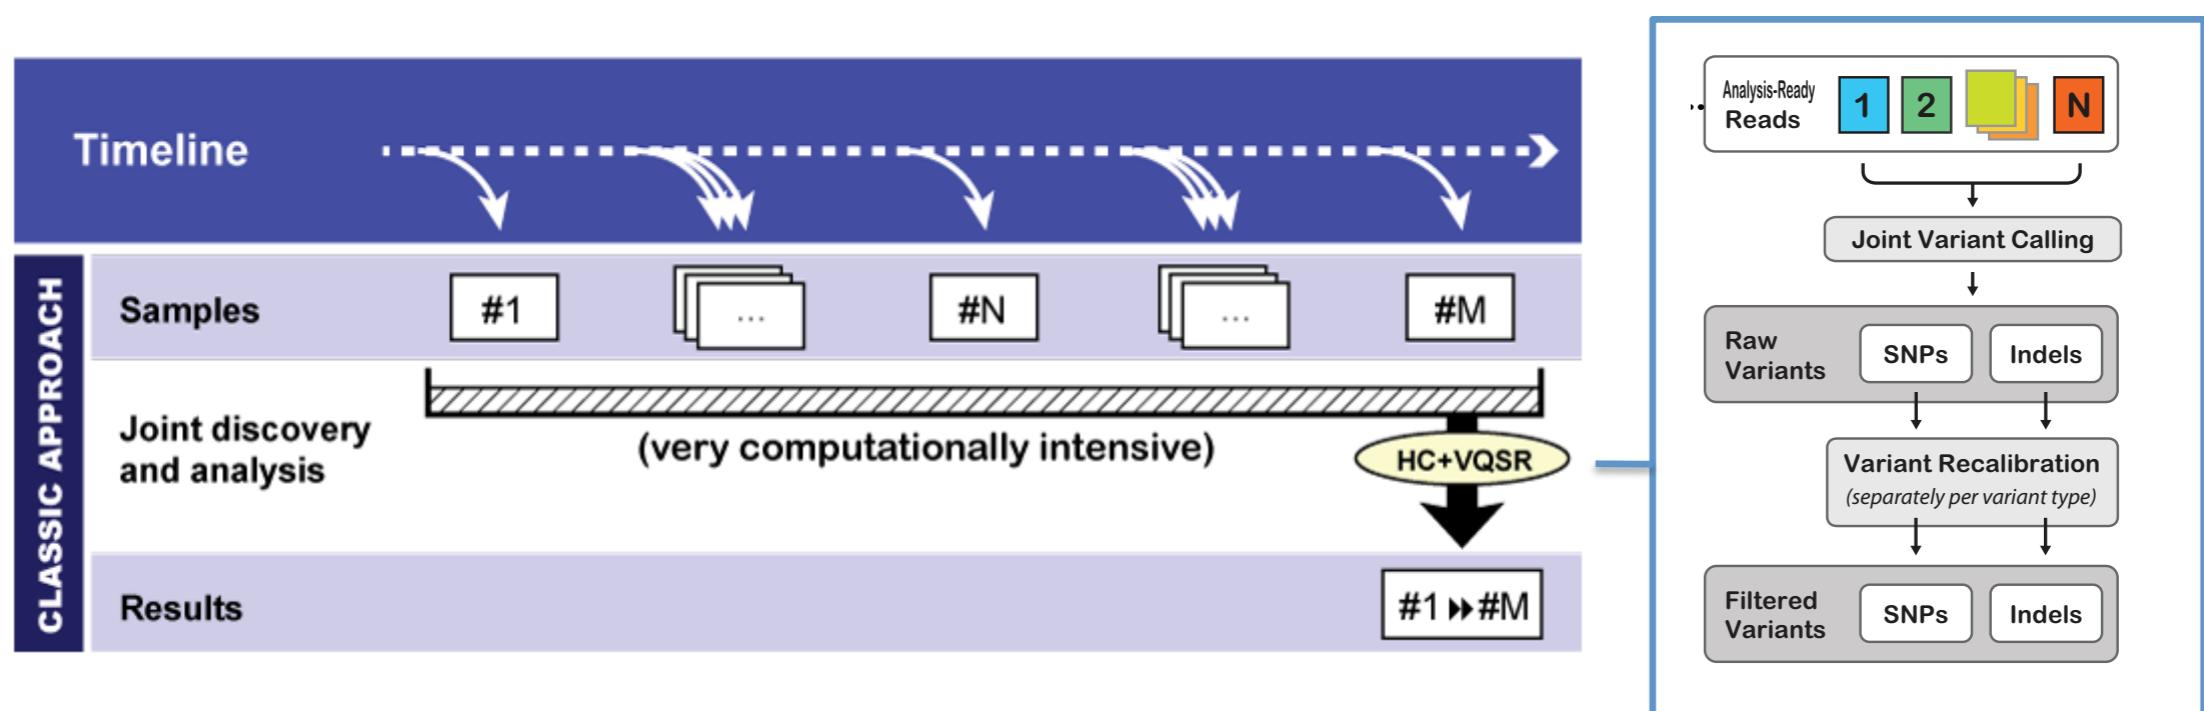
- Call variants jointly on all sample data
 - **Scales badly** -> limitations in amount of data that can be processed
 - Slow with **UnifiedGenotyper** (per-locus calculations)
 - *Impossibly* slow with **HaplotypeCaller** (so much extra work!)



But we want to use **HaplotypeCaller** because it is so much better!

Problems with the “all together” approach

- Computing costs
 - The “N+1 problem”



Regular* VCF

```
##fileformat  
##ALT  
##FILTER  
##FORMAT  
##INFO  
##contig  
##reference
```

HEADER

```
#record headers  
■ variant site record  
■ variant site record  
■ variant site record
```

RECORDS

* Some tools may output an all-sites VCF that looks like what you can get using HC with -ERC BP_RESOLUTION but they do not provide an accurate estimate of reference confidence.

HaplotypeCaller gVCF

-ERC GVCF

```
##fileformat  
##ALT  
##FILTER  
##FORMAT  
##GVCFBlock  
##INFO  
##contig  
##reference
```

-ERC BP_RESOLUTION

```
##fileformat  
##ALT  
##FILTER  
##FORMAT  
##INFO  
##contig  
##reference
```

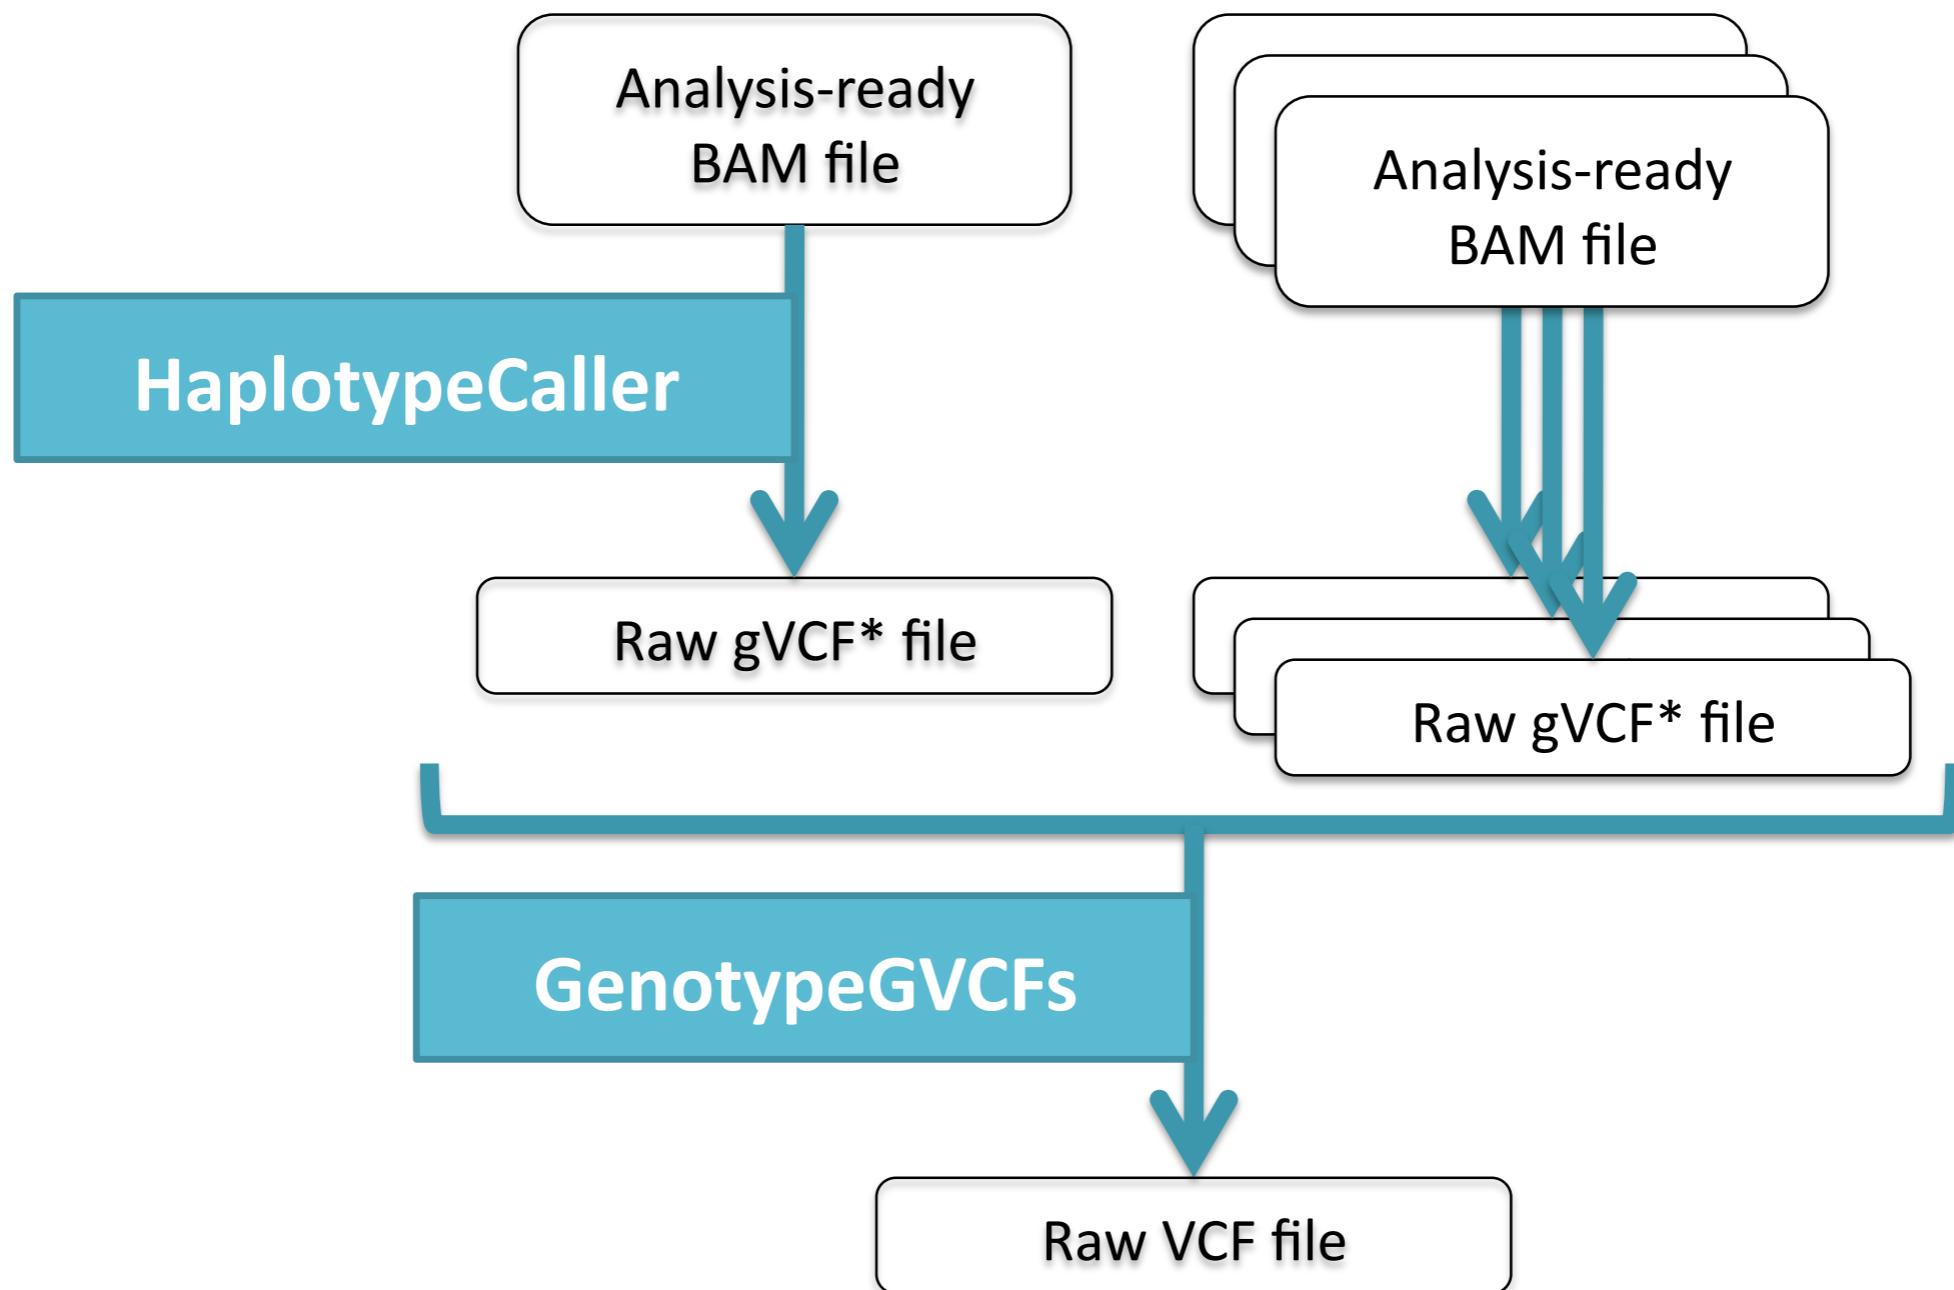
#record headers

```
■ non-variant block record  
■ variant site record  
■ non-variant block record  
■ variant site record  
■ non-variant block record  
■ variant site record  
■ non-variant block record
```

#record headers

```
■ non-variant site record  
■ variant site record  
■ non-variant site record  
■ non-variant site record  
■ non-variant site record  
■ variant site record  
■ non-variant site record  
■ non-variant site record  
■ variant site record  
■ non-variant site record  
■ non-variant site record  
■ non-variant site record
```

Variant calling + joint genotyping workflow



Variant annotations provide key information
to identify and remove artifacts!

VCF record for an A/G SNP at 22:49582364

22	49582364	.	A	G	198.96	.
AC=3;	AF=0.50;	AN=6;	DP=87;	MLEAC=3;	MLEAF=0.50;	MQ=71.31;
MQ0=22;	QD=2.29;	SB=-31.76	GT:DP:GQ	0/1:12:99.00	0/1:11:89.43	0/1:28:37.78
INFO field						
AC	No. chromosomes carrying alt allele	MLEAF	Max likelihood AF			
AN	Total no. of chromosomes	MQ	RMS MAPQ of all reads			
AF	Allele frequency	MQ0	No. of MAPQ 0 reads at locus			
DP	Depth of coverage	QD	QUAL score over depth			
MLEAC	Max likelihood AC	SB	Estimated strand bias score			

VCF Files store variant information

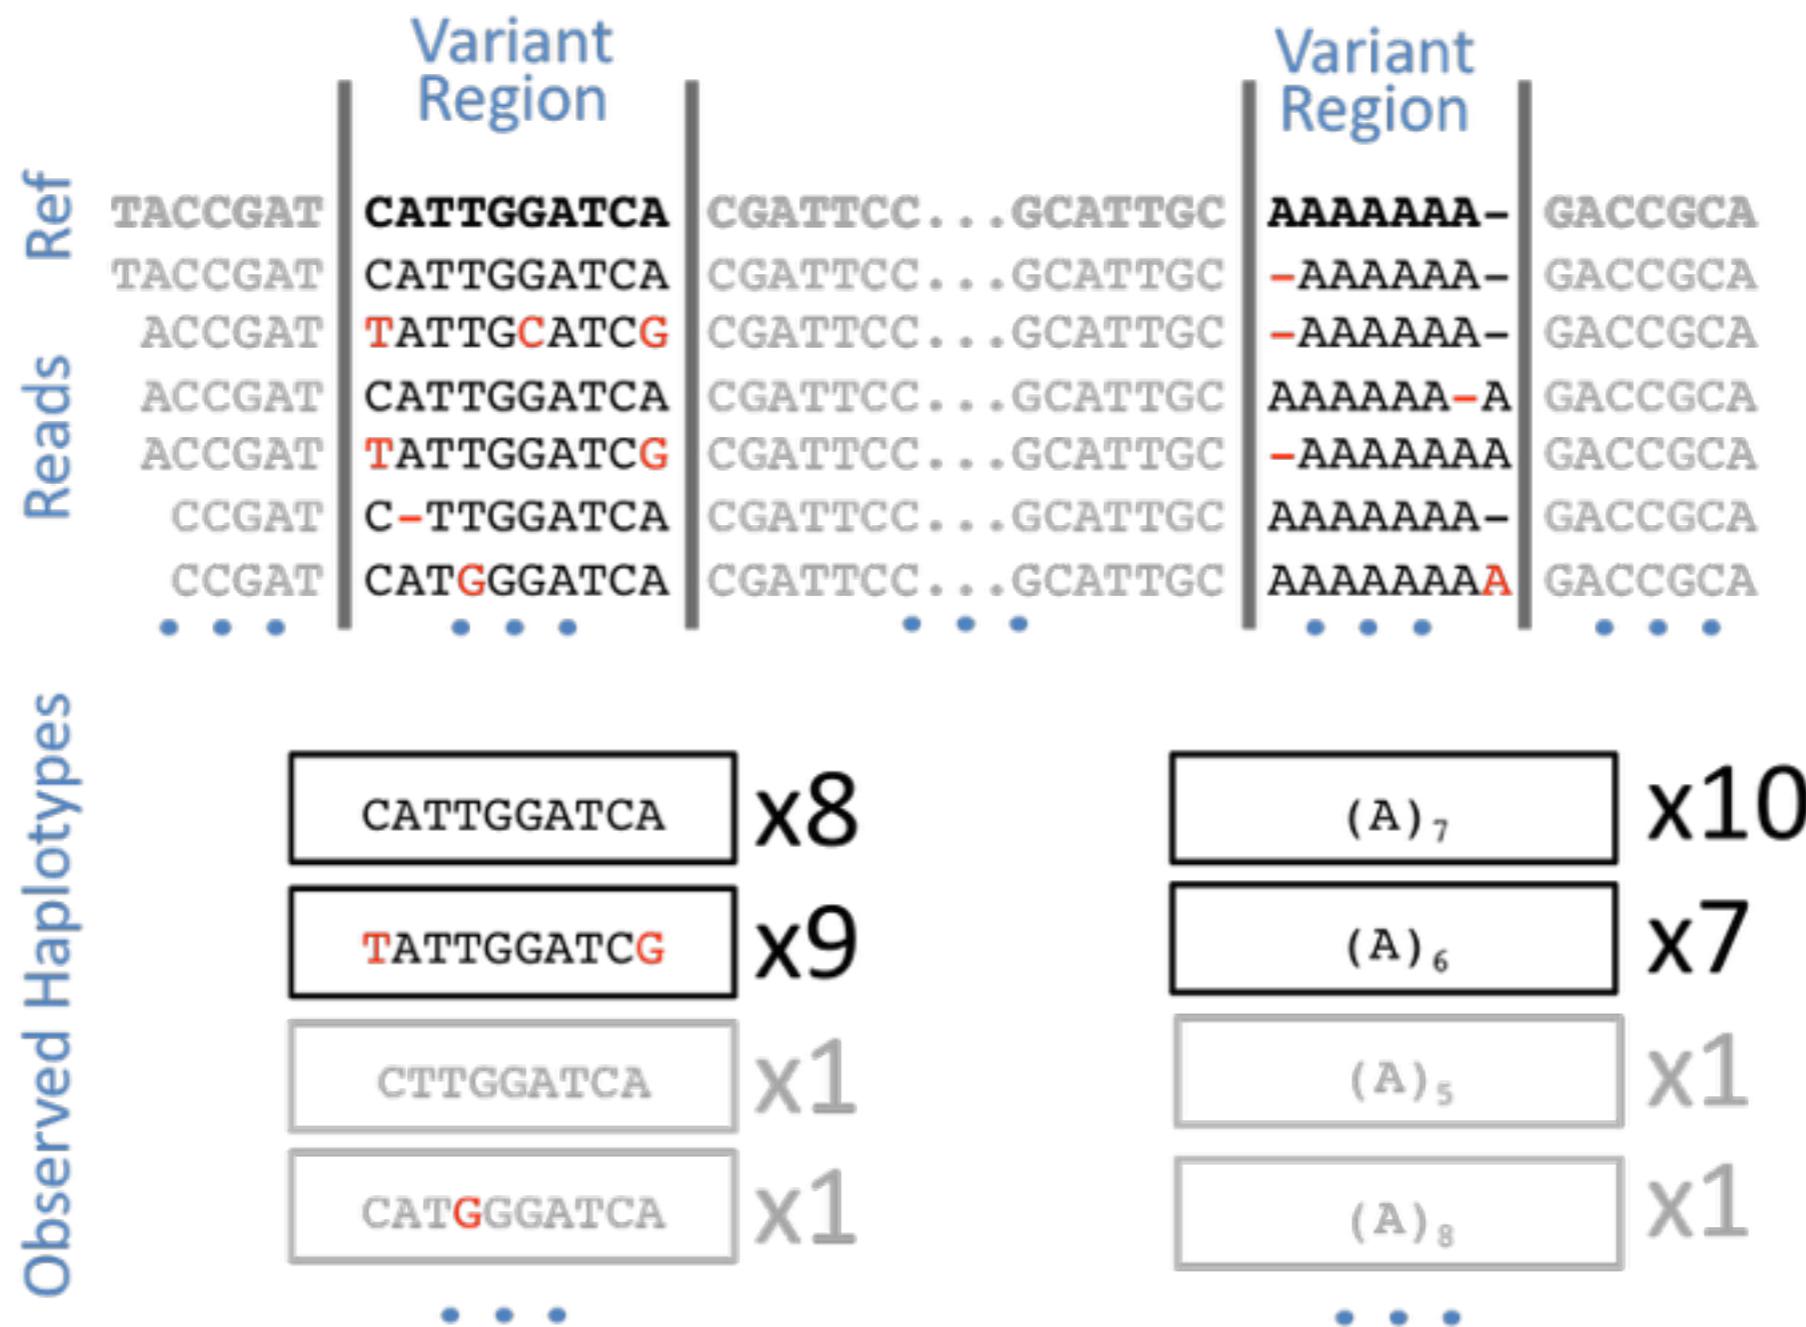
```
##fileformat=VCFv4.1
##reference=1000GenomesPilot-NCBI36
##INFO<ID=DP,Number=1>Type=Integer>Description="Total Depth">
##INFO<ID=AF,Number=A>Type=Float>Description="Allele Frequency">
##INFO<ID=DB,Number=0>Type=Flag>Description="dbSNP membership, build 129">
##FILTER<ID=s50>Description="Less than 50% of samples have data">
##FORMAT<ID=GT,Number=1>Type=String>Description="Genotype">
##FORMAT<ID=GQ,Number=1>Type=Integer>Description="Genotype Quality">
##FORMAT<ID=DP,Number=1>Type=Integer>Description="Read Depth">
#CHROM POS ID REF ALT QUAL FILTER INFO
#FORMAT FORMAT NA00001 NA00002 NA00003
20 14370 rs6054257 G A 29 PASS DP=14;AF=0.5;DB
    GT:GQ:DP 0|0:48:1 1|0:48:8 1/1:43:5
20 1110696 rs6040355 A G,T 67 PASS DP=10;AF=0.333,0.667;DB
    GT:GQ:DP 1|2:21:6 2|1:2:0 2/2:35:4
20 1230237 . T . 47 PASS DP=13
    GT:GQ:DP 0|0:54:7 0|0:48:4 0/0:61:2
20 1234567 microsat1 GTCT G,GTACT 50 PASS DP=9
    GT:GQ:DP 0/1:35:4 0/2:17:2 1/1:40:3
```

The diagram illustrates the structure of a VCF file. It features a red bracket on the right side that spans from the start of the header section to the end of the variant records section. The header section includes the file format, reference genome, and various metadata fields like INFO, FILTER, and FORMAT. The variant records section begins with a chromosome, position, and ID header, followed by a FORMAT header defining sample names (NA00001, NA00002, NA00003). Below this, several variants are listed, each with columns for chromosome, position, ID, reference allele, alternative allele, quality score, filter status, and detailed genotype information.

Official specification in

www.1000genomes.org/wiki/Analysis/Variant_Call_Format/vcf-variant-call-format-version-41

FreeBayes



FreeBayes

- Free and open source.
- Uses literal sequences of bases and haplotypes to call SNPs so is less affected by local alignment issues.
- Does not have “gvcf” n+1 method, although complicated work around exists.
- Generally faster than GATK, although RAM intensive.

Samtools mpileup

- Calls SNPs based on reads per site, no realignment.
- No longer being supported.